

Register No.

Date:

Lab Manual

NOSQL DATABASES LAB

III Year – II Semester VR -20

Name of the Student:

Regd.No.: Section:

Department:

Year: 20 -20

Jerripotula Priyanka, Asst Prof. CSE

VIGNAN'S INSTITUTE OF INFORMATION AND TECHNOLOGY,
DUVVADA

Register No.

Date:



VIGNAN's INSTITUTE OF INFORMATION TECHNOLOGY
(AUTONOMOUS)

(Approved by AICTE- New Delhi & Affiliated to JNTU-GV, Vizianagaram)
Beside VSEZ, Duvvada, Vadlapudi Post, Gajuwaka, Visakhapatnam - 530 049.

Department of
Computer Science and Engineering

CERTIFICATE

This is to certify that Mr./Ms. bearing
the Regd. No. is a student of B. Tech
Semester has successfully completed experiments in
..... Lab during the Academic Year 2023-2024.

Faculty In-charge

Head of the Department

External Examiner

Register No.

Date:

INDEX

Sl. No	Date	Title of the Experiment	Skill	Page No.	Marks	Sign. With Date
1		Exercise – 1 Introduction to MongoDB and its Installation on Windows & Linux	MongoDB Installation	1		
2		Exercise – 2 Description of mongo Shell, Create database and show database.	Database Creation	11		
3		Exercise – 3 Commands for MongoDB and to study operations in MongoDB – Insert, Query, Update, Delete and Projection	MongoDB Commands	13		
4		Exercise – 4 Download a zip code dataset at http://media.mongodb.org/zips.json . Use mongo import to import the zip code dataset into MongoDB. After importing the data, answer the following questions by using aggregation pipelines: 1. Find all the states that have a city called “BOSTON”. Find all the states and cities whose names include the string “BOST”. 2. Each city has several zip codes. Find the city in each state with the greatest number of zip codes and rank those cities along with the states using the city populations. MongoDB can query on spatial information	MongoDB	17		
5		Exercise – 5 Where Clause equivalent in MongoDB: Write a MongoDB query to find the restaurants that achieved a score, more than 80 but less than 100	MongoDB – where clause	21		
6		Exercise – 6 Write a MongoDB query to find the restaurants which locate in latitude value less than -95.754168.	MongoDB – Query	23		
7		Exercise – 7 To study operations in MongoDB – AND in MongoDB, OR in MongoDB, Limit Records and Sort Records. To study operations in MongoDB –	MongoDB Operations	25		

Register No.

Date:

		Indexing, Advanced Indexing, Aggregation and Map Reduce				
8		Exercise – 8 Column oriented databases study, queries, and practices	Column Oriented database	30		
9		Exercise – 9 Create a database that stores road cars. Cars have a manufacturer, a type. Each car has a maximum performance and a maximum torque value. Do the following: Test Cassandras replication schema and consistency models.	Cassandra	33		
10		Exercise – 10 Master Data Management using Neo4j Manage your master data more effectively the world of master data is changing. Data architects and application developers are swapping their relational databases with graph databases to store their master data. This switch enables them to use a data store optimized to discover new insights in existing data, provide a360-degree view of master data and answer questions about data relationships in real time	Data Management	45		
11		Exercise – 11 Shopping Mall case study using Cassandra, where we have many customers ordering items from thermall and, we have suppliers who deliver them their ordered items	Cassandra	48		

Register No.

Date:

Exercise-1

Introduction to MongoDB

MongoDB: Revolutionizing Data Management and Its Significance

In today's digital age, data is the lifeblood of businesses, fueling insights, innovation, and competitive advantage. However, traditional relational databases struggle to keep pace with the increasing volume, velocity, and variety of data generated daily. Enter MongoDB, a leading NoSQL database, heralded for its flexibility, scalability, and performance. This essay explores MongoDB's emergence, its key features, and its importance in modern data management.

MongoDB, founded in 2007 by Dwight Merriman, Eliot Horowitz, and Kevin Ryan, was born out of the need to address the limitations of relational databases in handling unstructured and semi-structured data. Unlike traditional SQL databases, MongoDB employs a document-oriented data model, storing data in flexible JSON-like documents. This schema-less approach allows for dynamic and agile data modeling, accommodating evolving application requirements without sacrificing performance.

One of MongoDB's core features is its scalability. With support for horizontal scaling through sharding, MongoDB can effortlessly distribute data across multiple servers, ensuring high availability and performance even as data volumes grow. This horizontal scaling capability is particularly crucial in today's era of big data, where organizations grapple with massive datasets generated by web applications, IoT devices, social media platforms, and more.

Another key aspect of MongoDB is its rich query language and powerful indexing capabilities. MongoDB's query language, which is based on JavaScript, provides developers with expressive and intuitive ways to retrieve and manipulate data. Furthermore, MongoDB's indexing capabilities enable efficient query execution, speeding up data retrieval and enhancing overall system performance. Combined, these features empower developers to build complex applications that deliver real-time insights and seamless user experiences.

Moreover, MongoDB's flexible data model makes it well-suited for agile development practices like DevOps and continuous integration/continuous deployment (CI/CD). Developers can quickly iterate on application features, adapting to changing business

requirements without cumbersome schema migrations or downtime. This agility fosters innovation and accelerates time-to-market, giving organizations a competitive edge in today's fast-paced digital landscape.

MongoDB's importance extends beyond its technical capabilities; it has become a cornerstone of modern data architectures, underpinning a wide array of applications across various industries. From e-commerce platforms to social media networks, from financial services to healthcare, MongoDB powers mission-critical applications that drive business growth and transformation. Its versatility and scalability make it an ideal choice for a broad spectrum of use cases, from content management and real-time analytics to personalization and recommendation engines.

Furthermore, MongoDB's vibrant ecosystem of tools and integrations enhances its appeal to developers and enterprises alike. From robust drivers for popular programming languages like Python, Java, and Node.js to comprehensive monitoring and management tools like MongoDB Atlas, the MongoDB ecosystem provides everything developers need to build, deploy, and manage MongoDB applications with ease.

Additionally, MongoDB's commitment to open source and community collaboration has fostered a thriving ecosystem of developers, contributors, and enthusiasts. The MongoDB community actively contributes to the platform's evolution, sharing best practices, developing plugins and extensions, and providing support and guidance through forums, meetups, and online communities. This collaborative spirit has been instrumental in MongoDB's success and continued innovation.

In conclusion, MongoDB has emerged as a game-changer in the realm of data management, offering unparalleled flexibility, scalability, and performance for modern applications. Its document-oriented data model, horizontal scalability, expressive query language, and vibrant ecosystem make it a preferred choice for developers and enterprises worldwide. As organizations strive to harness the power of data to drive innovation and growth, MongoDB stands as a testament to the transformative potential of NoSQL databases in the digital age.

Installation of MongoDB

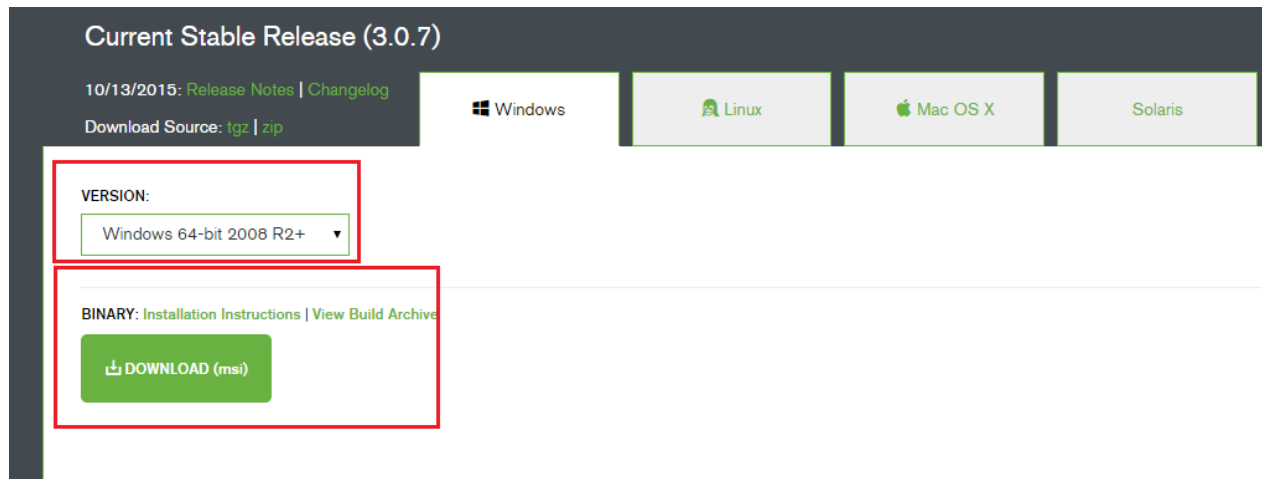
Mongo DB installation document for Windows OS: Below installation is done on Windows 10:

1. Download msi file from below link:

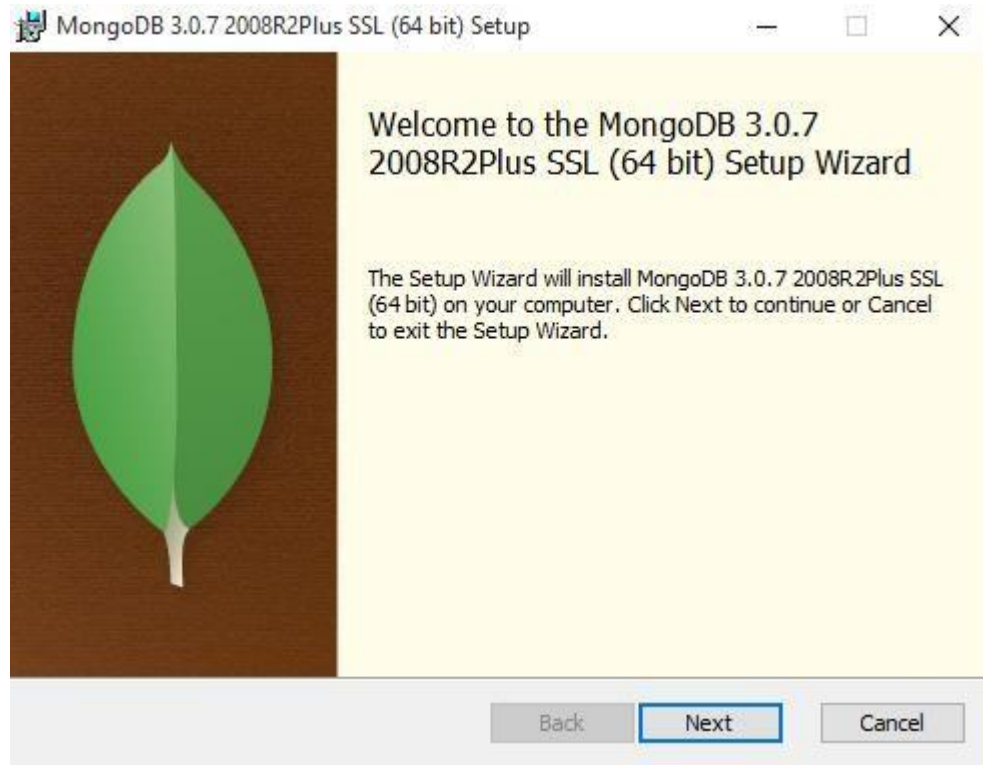
Register No.

Date:

<https://www.mongodb.org/downloads#production>



2. Run the msi file

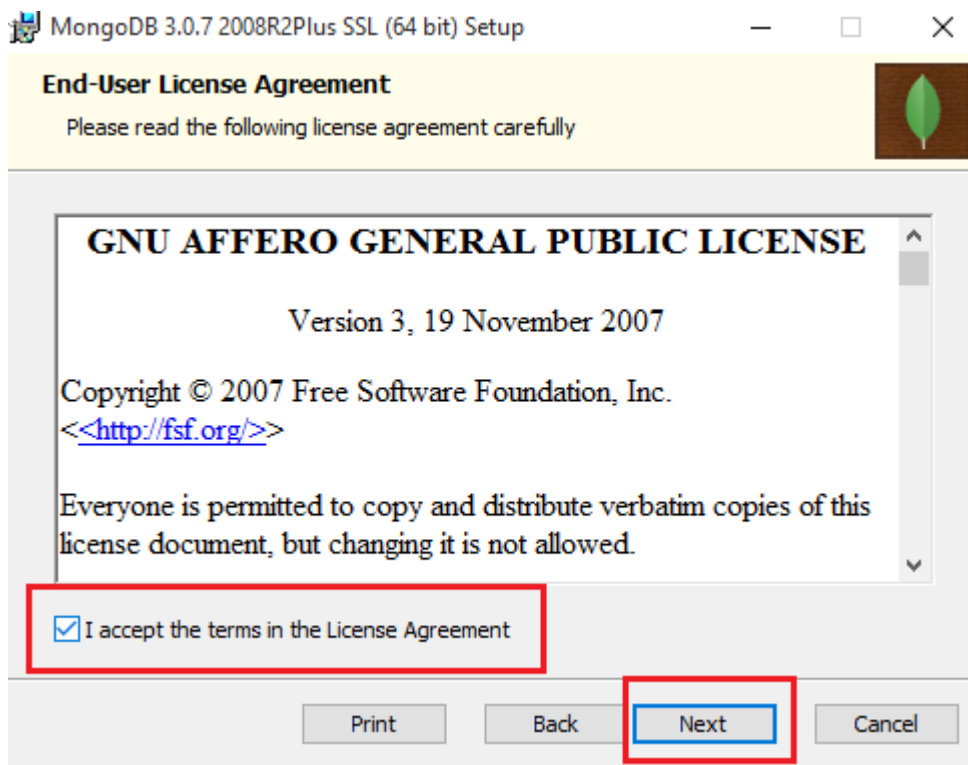


Click Next

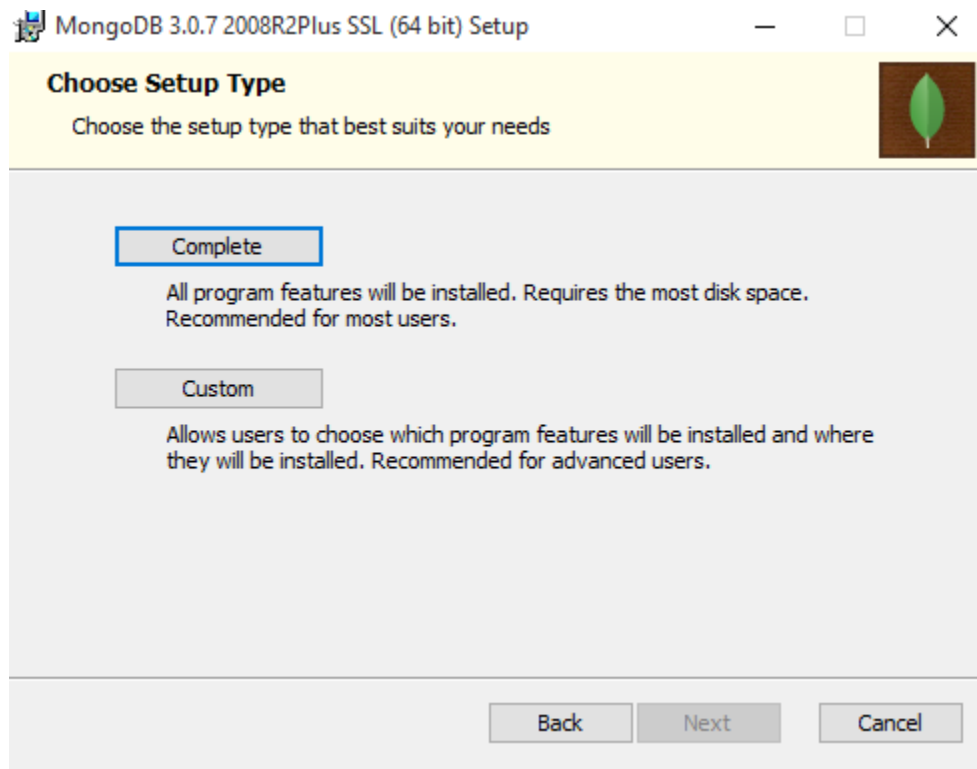
3. Accept the terms in the License Agreement and select Next

Register No.....

Date:.....




4. Select Complete




5. Select Install

Register No.....

Date:.....




6. Wait until you get the finish screen

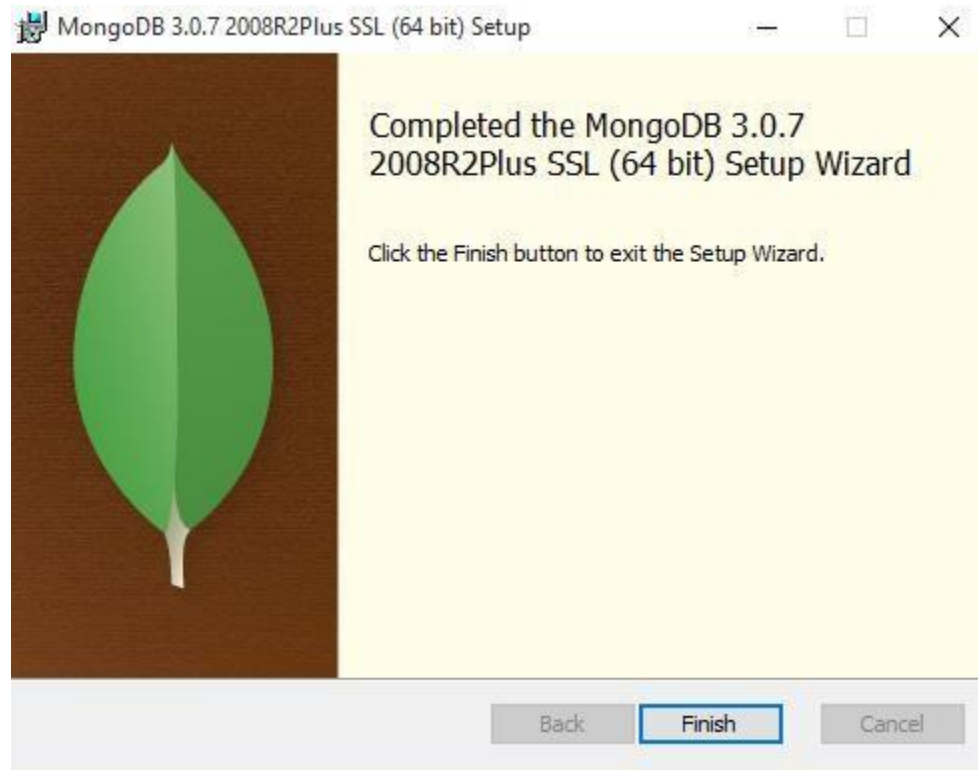


Register No.....

Date:.....




7. Click Finish




Register No.....

Date:.....

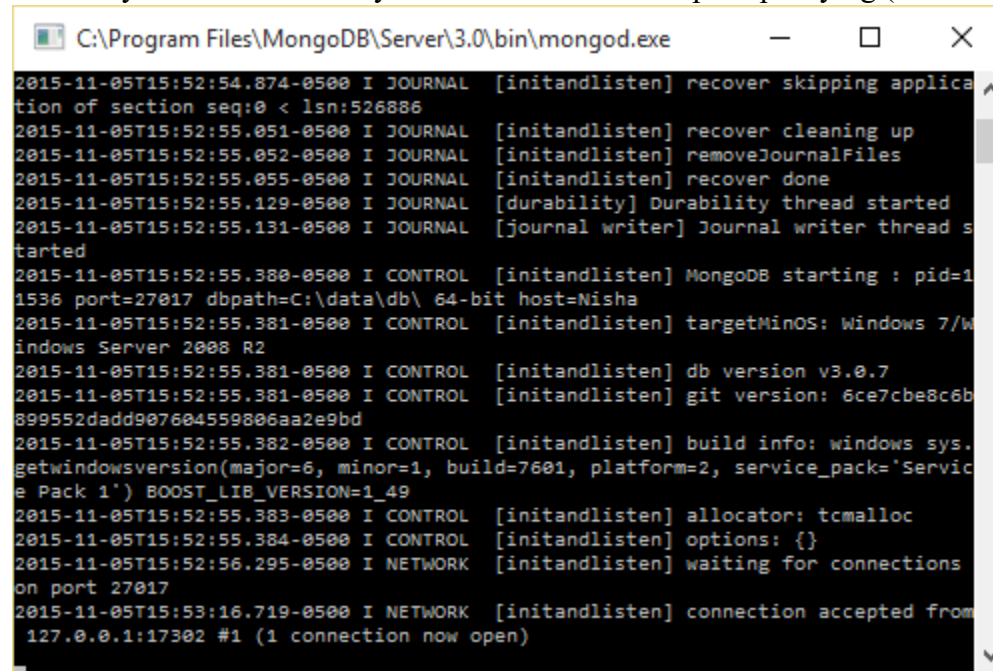
8. Create folders in C drive: data>db
9. Add below path in the environment variable
10. Navigate to C:> ProgramFiles>MongoDB>Server>3.0>bin



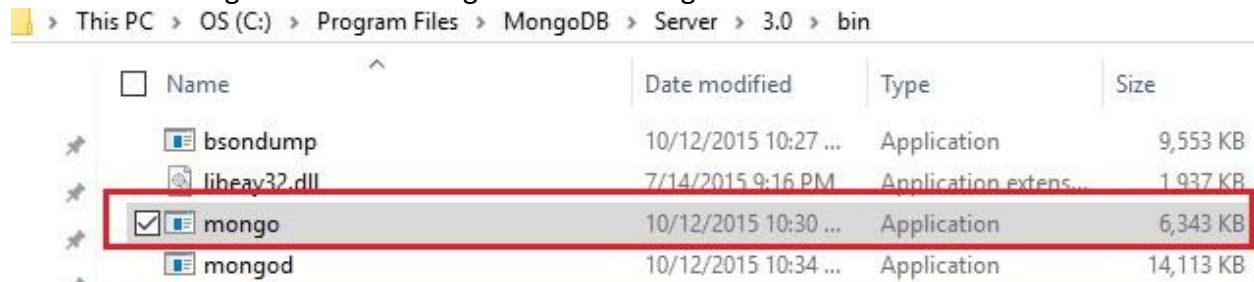
11. Right click and run as administrator: Start the mongo db server




Once you start the server you will see a command prompt saying (1 connection now open)



12. Start the mongo db database: Right Click on Mongo and run it as administration



Once your database is started, it will be connected to the test database by default. Use command: `show dbs` → it will show you all dbs



MongoDB stores data in the form of Documents, which are JSON-like field –value pairs. Mongo DB Documents are called BSON documents: (Binary representation of JSON) with additional type information.

```

{
  name: "sue",
  age: 26,
  status: "A",
  groups: [ "news", "sports" ]
}
  
```

← field: value
 ← field: value
 ← field: value
 ← field: value

MongoDB stores all documents in [collections](#). A collection is a group of related documents that have a set of shared common indexes. Collections are analogous to a table in relational databases.



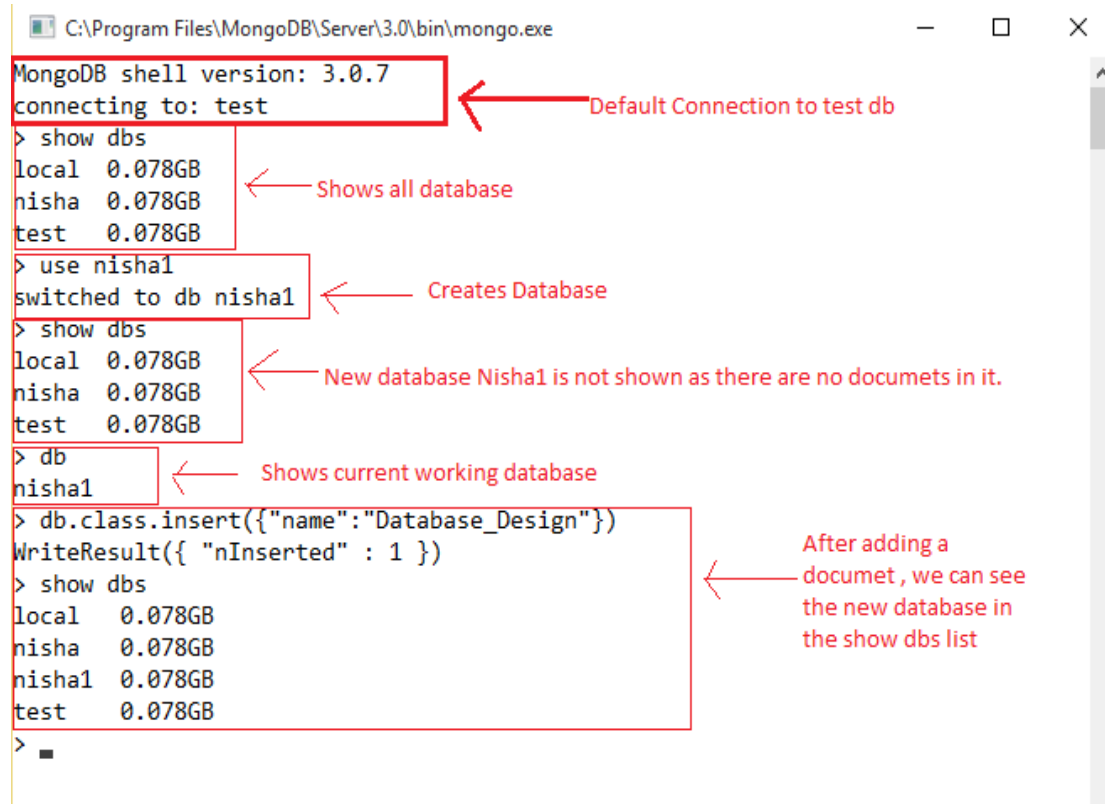
Use dbname → created database e.g. use nisha1

```
C:\Program Files\MongoDB\Server\3.0\bin>mongo.exe
MongoDB shell version: 3.0.7
connecting to: test
> show dbs
local  0.078GB
nisha  0.078GB
test   0.078GB
> use nisha1
switched to db nisha1
>
```

Good resource for learning Mongo DB:

http://www.tutorialspoint.com/mongodb/mongodb_create_collection.htm

Some basic commands to try:



The screenshot shows a Windows command prompt window titled "C:\Program Files\MongoDB\Server\3.0\bin\mongo.exe". The terminal output is as follows:

```
MongoDB shell version: 3.0.7
connecting to: test
> show dbs
local  0.078GB
nisha  0.078GB
test   0.078GB
> use nisha1
switched to db nisha1
> show dbs
local  0.078GB
nisha  0.078GB
test   0.078GB
> db
nisha1
> db.class.insert({"name":"Database_Design"})
WriteResult({ "nInserted" : 1 })
> show dbs
local  0.078GB
nisha  0.078GB
nisha1 0.078GB
test   0.078GB
> █
```

Annotations with red arrows point to specific parts of the output:

- "Default Connection to test db" points to "connecting to: test".
- "Shows all database" points to the output of the first "show dbs" command.
- "Creates Database" points to "switched to db nisha1".
- "New database Nisha1 is not shown as there are no documets in it." points to the second "show dbs" output.
- "Shows current working database" points to the output of the "db" command.
- "After adding a documet , we can see the new database in the show dbs list" points to the third "show dbs" output, which now includes "nisha1".

References: <https://docs.mongodb.org/manual/>

Exercise-2

Description of mongo Shell, Create database and show database

Navigating MongoDB: Unveiling the MongoDB Shell's Power in Creating and Managing Databases

In the realm of modern data management, MongoDB stands out as a frontrunner, offering unparalleled flexibility and scalability. At the heart of MongoDB's ecosystem lies the MongoDB Shell, a command-line interface (CLI) that empowers developers and administrators to interact with MongoDB instances seamlessly. This essay embarks on an exploration of the MongoDB Shell, elucidating its functionalities, and demonstrating the process of creating and managing databases with illustrative examples.

Understanding the MongoDB Shell:

The MongoDB Shell, often referred to as "mongo," serves as a robust tool for interfacing with MongoDB databases directly from the command line. Powered by JavaScript, the MongoDB Shell provides a familiar environment for executing database operations, administering collections, and querying data effortlessly. It is an indispensable component of MongoDB's toolkit, offering both simplicity and power to users.

Creating Databases with MongoDB Shell:

MongoDB's schema-less architecture allows for dynamic database creation, enabling developers to effortlessly set up databases on-the-fly. The MongoDB Shell simplifies this process, providing a straightforward mechanism for creating databases. Let's delve into a practical example:

```
// Start MongoDB Shell
```

```
mongo
```

```
// Create a new database named "mydatabase"
```

```
use mydatabase
```

In the above example, we initiate the MongoDB Shell by running the ``mongo`` command in the terminal. Once inside the shell, the ``use`` command is employed to switch to the desired database or create it if it does not already exist. Upon execution, MongoDB switches to the "mydatabase" database, creating it if it's not present. This lazy creation approach ensures database creation only upon necessity, conserving resources efficiently.

Exploring Databases:

After creating a database, users often need to navigate and explore existing databases. The MongoDB Shell provides commands to facilitate this exploration, allowing users to view available databases effortlessly. Let's illustrate this with an example:

```
// Display a list of all databases
```

```
show databases
```

Executing the ``show databases`` command within the MongoDB Shell yields a comprehensive list of databases available on the MongoDB server. The output includes system databases such as "admin" and "local," along with any user-defined databases like "mydatabase." This command provides users with a quick overview of the databases present, aiding in navigation and administration tasks.

Managing Collections:

Within MongoDB databases, data is organized into collections, analogous to tables in relational databases. The MongoDB Shell enables users to manage collections efficiently, facilitating tasks such as viewing existing collections and creating new ones. Let's illustrate this with an example:

```
// Display a list of collections in the current database
```

```
show collections
```

Executing the ``show collections`` command within a database reveals a list of all collections present within that database. Collections play a pivotal role in MongoDB's data model, serving as containers for JSON-like documents. This command empowers users to quickly ascertain the structure of their database, aiding in data exploration and analysis.

Conclusion:

In conclusion, the MongoDB Shell serves as a cornerstone of MongoDB's ecosystem, providing developers and administrators with a powerful interface for managing databases effortlessly. Through simple commands like ``use``, ``show databases``, and ``show collections``, users can create databases, navigate existing ones, and administer collections seamlessly. With its intuitive JavaScript-based interface, the MongoDB Shell democratizes database management, empowering users of all levels to harness the full potential of MongoDB's flexible and scalable architecture.

By bridging the gap between developers and databases, the MongoDB Shell fosters collaboration, innovation, and efficiency, driving forward the evolution of modern data management practices. As organizations continue to embrace MongoDB as their preferred database solution, the MongoDB Shell will remain an indispensable tool in their arsenal, enabling them to navigate the complexities of data management with confidence and ease.

Exercise-3

Commands for MongoDB and to study operations in MongoDB – Insert, Query, Update, Delete and Projection

MongoDB Commands: Exploring Operations

MongoDB, a leading NoSQL database, offers a rich set of commands for performing various operations, from basic CRUD (Create, Read, Update, Delete) tasks to more advanced operations like Projection. In this essay, we delve into five fundamental MongoDB commands, namely Insert, Query, Update, Delete, and Projection, exploring their functionalities and providing illustrative examples with output.

1. Insert:

The `insert` command in MongoDB is used to insert documents into a collection. Documents in MongoDB are JSON-like objects that store data in key-value pairs. Let's explore the `insert` command with an example:

```
// Insert a document into the "users" collection
```

```
db.users.insert({  
  name: "John Doe",  
  age: 30,  
  email: "john@example.com"  
})
```

In the above example, we insert a document representing a user into the "users" collection. The document contains fields such as "name," "age," and "email," each with corresponding values. Upon execution, MongoDB adds this document to the "users" collection.

****Output:****

```
WriteResult({ "nInserted" : 1 })
```

The output indicates that one document was successfully inserted into the collection.

2. Query:

The `find` command in MongoDB is used to query documents from a collection based on specified criteria. It allows users to retrieve data that matches certain conditions. Let's explore the `find` command with an example:

```
// Query documents where age is greater than 25
```

```
db.users.find({ age: { $gt: 25 } })
```

In the above example, we query the "users" collection to retrieve documents where the "age" field is greater than 25. MongoDB's query language supports various operators like `$gt` (greater than), `$lt` (less than), `$eq` (equal to), etc., enabling users to craft complex queries.

****Output:****

```
{ "_id" : ObjectId("60b7d56da8e1d54b08a789cf"), "name" : "John Doe", "age" : 30, "email" :  
"john@example.com" }
```

The output displays the document(s) that match the specified criteria.

3. Update:

The `update` command in MongoDB is used to modify existing documents within a collection. It allows users to update specific fields or replace entire documents. Let's explore the `update` command with an example:

```
// Update the email of a user with a specific name
```

```
db.users.update(  
  { name: "John Doe" },  
  { $set: { email: "john.doe@example.com" } }  
)
```

In the above example, we update the email address of the user with the name "John Doe" to "john.doe@example.com." MongoDB's update command supports various modifiers like `$set`, `$unset`, `$inc`, etc., enabling precise updates to documents.

****Output:****

```
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
```

The output indicates that one document was matched and modified as a result of the update operation.

4. Delete:

The ``delete`` command in MongoDB is used to remove documents from a collection that match specified criteria. It allows users to delete one or multiple documents based on the provided query. Let's explore the ``delete`` command with an example:

```
// Delete documents where age is less than 25
```

```
db.users.deleteMany({ age: { $lt: 25 } })
```

In the above example, we delete documents from the "users" collection where the "age" field is less than 25. MongoDB's ``deleteMany`` command removes all documents that match the specified criteria.

****Output:****

```
{ "acknowledged" : true, "deletedCount" : 0 }
```

The output indicates that no documents were deleted as none matched the specified criteria.

5. Projection:

The ``find`` command in MongoDB also supports Projection, allowing users to retrieve only specific fields from documents rather than the entire document. It helps in optimizing query performance and reducing network overhead. Let's explore Projection with an example:

```
// Retrieve only the name and email fields of users
```

```
db.users.find({}, { name: 1, email: 1, _id: 0 })
```

In the above example, we query the "users" collection to retrieve only the "name" and "email" fields of documents while excluding the "_id" field. This reduces the amount of data transferred over the network and improves query performance.

****Output:****

```
{ "name" : "John Doe", "email" : "john.doe@example.com" }
```

The output displays the specified fields of the document(s) retrieved from the collection.

Conclusion:

In conclusion, MongoDB offers a versatile set of commands for performing various operations on databases and collections. From inserting and querying data to updating, deleting, and projecting fields, MongoDB commands provide users with the flexibility and power to manipulate data effectively. Through illustrative examples and output, this essay elucidates the functionalities of MongoDB commands, empowering users to leverage MongoDB's capabilities for efficient data management and manipulation.

Exercise-4

Download a zip code dataset at <http://media.mongodb.org/zip.js>. Use mongo import to import the zip code dataset into MongoDB. After importing the data, answer the following questions by using aggregation pipelines:

1. Find all the states that have a city called "BOSTON". Find all the states and cities whose names include the string "BOST".
2. Each city has several zip codes. Find the city in each state with the greatest number of zip codes and rank those cities along with the states using the city populations. MongoDB can query on spatial information

Aim: Download a zip code dataset at <http://media.mongodb.org/zip.js>. Use mongo import to import the zip code dataset into MongoDB. After importing the data, answer the following questions by using aggregation pipelines:

1. Find all the states that have a city called "BOSTON". Find all the states and cities whose names include the string "BOST".
2. Each city has several zip codes. Find the city in each state with the greatest number of zip codes and rank those cities along with the states using the city populations. MongoDB can query on spatial information.

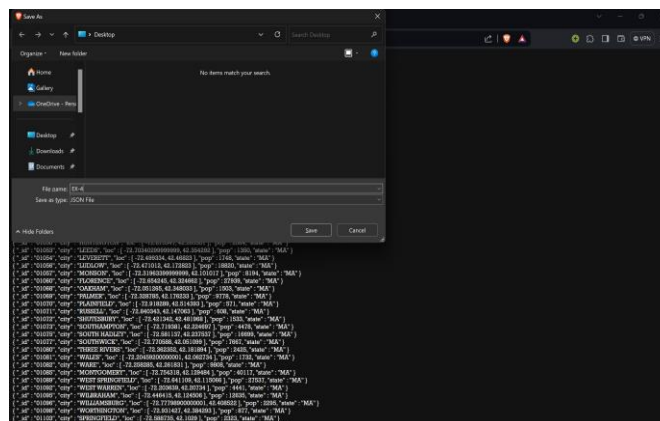
PROCEDURE

STEPS

1. **Downloading the Dataset.**

Click on the link below.

<http://media.mongodb.org/zip.js>



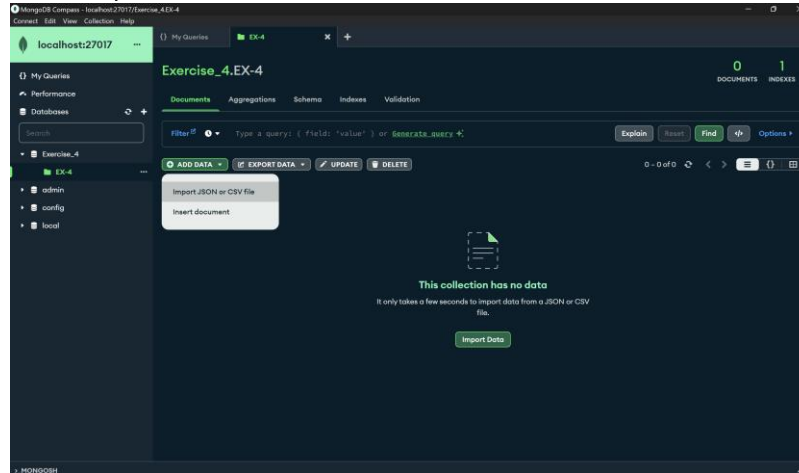
Press '**Ctrl + S**' to save the file.

File Name: EX-4

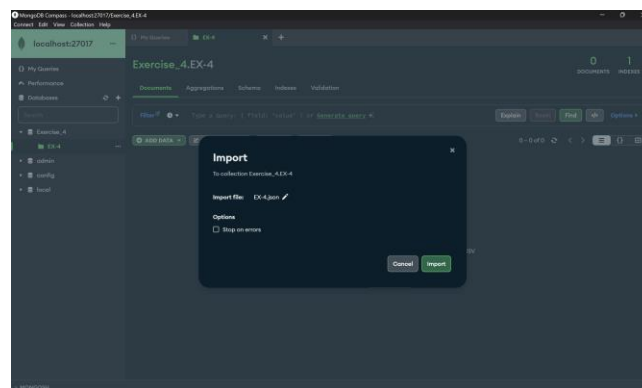
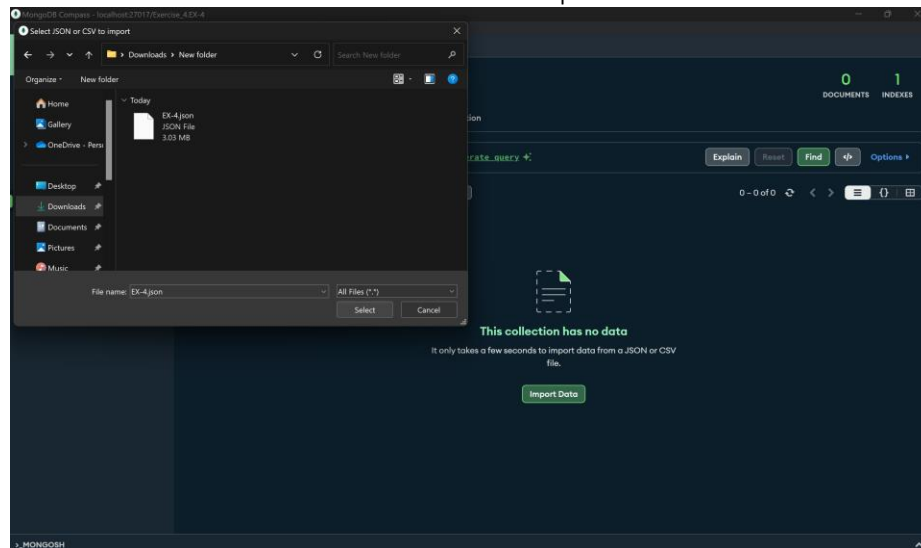
File Type: JSON File

2. Import the Dataset on to MongoDB Compass

- Connect to MongoDB Compass
- Create database. For example, LAB.
- Create a collection. For example, EX4.
- Click on **ADD DATA**. Select **Import JSON or SCV File**.



- Chose the download dataset. And click **Select**. And then Import.



- Open MongoDB shell or **_MONGOSH**.

3. Queries

a) Commands

```
test> use LAB
```

```
LAB> db.EX4.distinct("state", { "city": "BOSTON" })
```

```
LAB> db.EX4.find({"$or": [{"state": {"$regex": ".*BOST.*"}}, {"city": {"$regex": ".*BOST.*"}}], {"state": 1, "city": 1, "_id": 0});
```

The screenshot shows the MongoDB Compass application. The top bar indicates the connection to 'localhost:27017/LAB.EX4'. The left sidebar shows the 'My Queries' tab with a query named 'LAB.EX4'. The main panel displays the MongoDB shell commands and their results. The first command is 'use LAB', which switches the database to 'LAB'. The second command is 'db.EX4.distinct("state", { "city": "BOSTON" })', which returns an array of state abbreviations: ['GA', 'KY', 'MA', 'NY', 'PA', 'TX', 'VA']. The third command is 'db.EX4.find({"\$or": [{"state": {"\$regex": ".*BOST.*"}}, {"city": {"\$regex": ".*BOST.*"}}], {"state": 1, "city": 1, "_id": 0});', which returns an array of documents where the state is 'BOSTON' and the city is 'MA'.

```

> use LAB
switched to db LAB
> db.EX4.distinct("state", { "city": "BOSTON" })
[
  'GA', 'KY',
  'MA', 'NY',
  'PA', 'TX',
  'VA'
]
> db.EX4.find({"$or": [{"state": {"$regex": ".*BOST.*"}}, {"city": {"$regex": ".*BOST.*"}}], {"state": 1, "city": 1, "_id": 0});
[
  {
    city: 'BOSTON',
    state: 'MA'
  },
  {
    city: 'BOSTON',
    state: 'MA'
  },
  {
    city: 'BOSTON',
    state: 'MA'
  },
  {
    city: 'BOSTON',
    state: 'MA'
  },
  {
    city: 'BOSTON',
    state: 'MA'
  }
]

```

b) Commands

```
LAB> cls
```

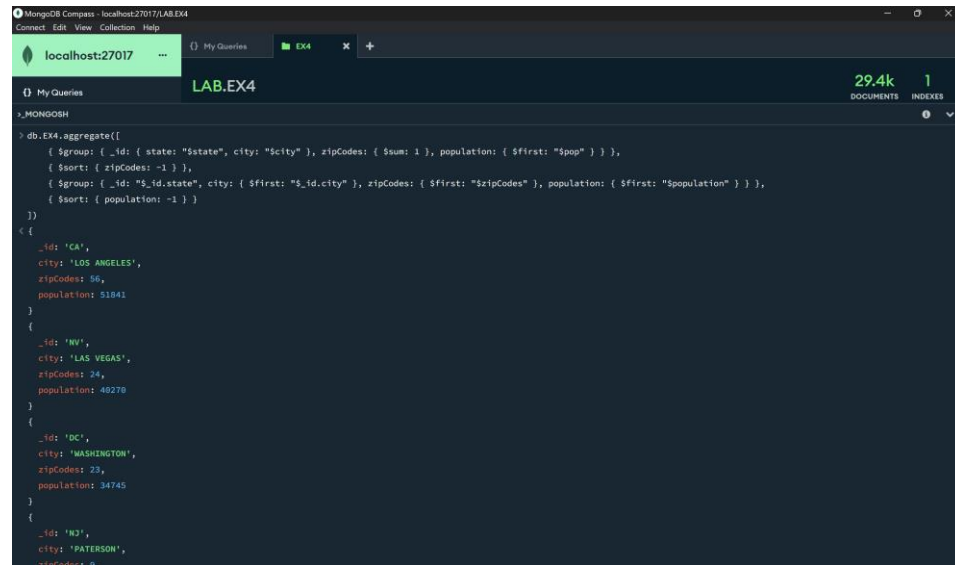
```

LAB> db.EX4.aggregate([
  { $group: { _id: { state: "$state", city: "$city" }, zipCodes: { $sum: 1 }, pop: { $first: "$pop" } } },
  { $sort: { zipCodes:-1 } },
  { $group: { _id: "$_id.state", city: { $first: "$_id.city" }, zipCodes: { $first: "$zipCodes" }, pop: { $first:
"$pop" } } },
  { $sort: { pop:-1 } }
])

```


Register No.

Date:



The screenshot shows the MongoDB Compass interface. The top bar indicates the connection to 'localhost:27017/LAB.EX4'. The left sidebar shows the 'My Queries' tab with a query named 'LAB.EX4'. The main panel displays the following aggregation query:

```
> db.EX4.aggregate([
  { $group: { _id: { state: "$state", city: "$city" }, zipCodes: { $sum: 1 }, population: { $first: "$pop" } } },
  { $sort: { zipCodes: -1 } },
  { $group: { _id: "$_id.state", city: { $first: "$_id.city" }, zipCodes: { $first: "$zipCodes" }, population: { $first: "$population" } } },
  { $sort: { population: -1 } }
])
```

The results are displayed in a table with the following columns: _id, city, zipCodes, and population. The results are sorted by population in descending order.

_id	city	zipCodes	population
CA	LOS ANGELES	56	51841
NV	LAS VEGAS	24	48278
DC	WASHINGTON	23	34745
ND	PATERSON	9	

Exercise-5

Where Clause equivalent in MongoDB: Write a MongoDB query to find the restaurants that achieved a score, more than 80 but less than 100

AIM:

Where Clause equivalent in MongoDB: Write a MongoDB query to find the restaurants that achieved a score, more than 80 but less than 100.

PROCEDURE

STEPS

1. MongoDB Compass

- i. Connect to MongoDB Compass
- ii. Create a collection. For example, EX5.

2. Commands

```
LAB> db.EX5.insertOne({name: "McDonalds", score: 85})
```

```
db.EX5.insertOne({name: "Dominos", score: 70})
```

```
db.EX5.insertOne({name: "Pizza Hut", score: 95})
```

```
db.EX5.insertOne({name: "Pastry Chef", score: 70})
```

```
db.EX5.insertOne({name: "Burger King", score: 88})
```

```
LAB> let newDocuments = [
```

```
  {name: "KFC", score: 92},
```

```
  {name: "Subway", score: 79},
```

```
  {name: "Starbucks", score: 91},
```

```
  {name: "Dunkin Donuts", score: 77},
```

```
  {name: "Taco Bell", score: 93}
```

```
]
```

```
db.EX5.insertMany(newDocuments)
```

```
LAB> db.EX5.find({ score: { $gt: 80, $lt: 100 } })
```

Register No.

Date: _____

```

# MongoDB Compass: localhost:27017/foodies
Connect Edit View Collections Tools

localhost:27017
My Queries
LAB.EXS
0
10000000000
10000000000

1 MongoDB Compass

db.foodies.find( { score: { $gt: 80, $lt: 100 } } )

{
  "_id": ObjectId("61579a5a679326a07f304337f1"),
  "name": "McDonalds",
  "score": 85
}

{
  "_id": ObjectId("61579a5a679326a07f304337f2"),
  "name": "Pizza Hut",
  "score": 85
}

{
  "_id": ObjectId("61579a5a679326a07f304337f3"),
  "name": "Burger King",
  "score": 86
}

{
  "_id": ObjectId("61579a5a679326a07f304337f4"),
  "name": "KFC",
  "score": 82
}

{
  "_id": ObjectId("61579a5a679326a07f304337f5"),
  "name": "Starbucks",
  "score": 91
}

}

[tab]

```

Exercise-6

Write a MongoDB query to find the restaurants which locate in longitude value less than -95.754168.

AIM

Write a MongoDB query to find the restaurants which locate in longitude value less than -95.754168.

PROCEDURE

STEPS

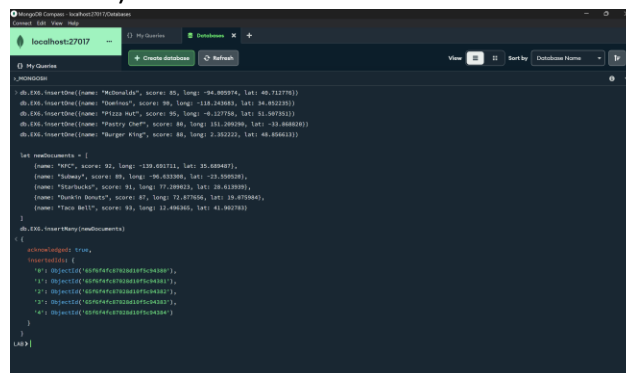
1. MongoDB Compass

- i. Connect to MongoDB Compass
- ii. Create a collection. For example, EX6.

2. Commands

```
LAB>db.EX6.insertOne({name: "McDonalds", score: 85, long: -94.005974, lat: 40.712776})
db.EX6.insertOne({name: "Dominos", score: 90, long: -118.243683, lat: 34.052235})
db.EX6.insertOne({name: "Pizza Hut", score: 95, long: -0.127758, lat: 51.507351})
db.EX6.insertOne({name: "Pastry Chef", score: 80, long: 151.209290, lat: -33.868820})
db.EX6.insertOne({name: "Burger King", score: 88, long: 2.352222, lat: 48.856613})
```

```
let newDocuments = [
  {name: "KFC", score: 92, long: -139.691711, lat: 35.689487},
  {name: "Subway", score: 89, long: -96.633308, lat: -23.550520},
  {name: "Starbucks", score: 91, long: 77.209023, lat: 28.613939},
  {name: "Dunkin Donuts", score: 87, long: 72.877656, lat: 19.075984},
  {name: "Taco Bell", score: 93, long: 12.496365, lat: 41.902783}
]
db.EX6.insertMany(newDocuments)
```



```

MongoDB Compass - localhost:27017
localhost:27017
My Databases
My Collections
View
Sort By
Database Name
db.EX6.insertOne({name: "McDonalds", score: 85, long: -94.005974, lat: 40.712776})
db.EX6.insertOne({name: "Dominos", score: 90, long: -118.243683, lat: 34.052235})
db.EX6.insertOne({name: "Pizza Hut", score: 95, long: -0.127758, lat: 51.507351})
db.EX6.insertOne({name: "Pastry Chef", score: 80, long: 151.209290, lat: -33.868820})
db.EX6.insertOne({name: "Burger King", score: 88, long: 2.352222, lat: 48.856613})

let newDocuments = [
  {name: "KFC", score: 92, long: -139.691711, lat: 35.689487},
  {name: "Subway", score: 89, long: -96.633308, lat: -23.550520},
  {name: "Starbucks", score: 91, long: 77.209023, lat: 28.613939},
  {name: "Dunkin Donuts", score: 87, long: 72.877656, lat: 19.075984},
  {name: "Taco Bell", score: 93, long: 12.496365, lat: 41.902783}
]
db.EX6.insertMany(newDocuments)

{
  "acknowledged": true,
  "insertedCount": 5,
  "insertedIds": [
    "60f0f0f0f0f0f0f0f0f0f0f0",
    "60f0f0f0f0f0f0f0f0f0f0f0",
    "60f0f0f0f0f0f0f0f0f0f0f0",
    "60f0f0f0f0f0f0f0f0f0f0f0",
    "60f0f0f0f0f0f0f0f0f0f0f0"
  ]
}
```

```
LAB> db.EX6.find({ long: { $lt: -95.754168 } })
```

Register No.

Date: _____

[illegible]

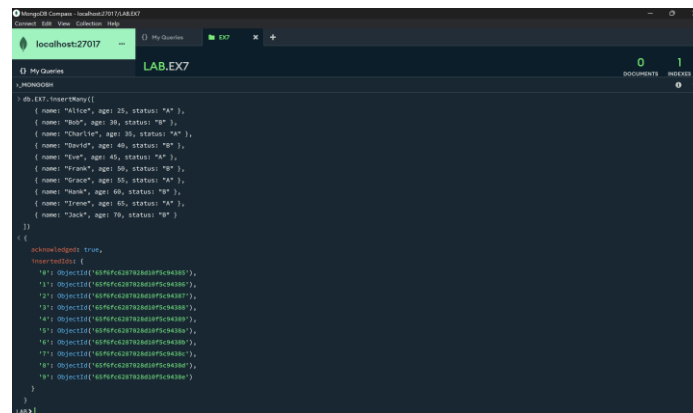
Exercise-7

AIM

To study operations in MongoDB – AND in MongoDB, OR in MongoDB, Limit Records and Sort Records. To study operations in MongoDB – Indexing, Advanced Indexing, Aggregation and Map Reduce.

DATABASE

```
LAB>db.EX7.insertMany([
  { name: "Alice", age: 25, status: "A" },
  { name: "Bob", age: 30, status: "B" },
  { name: "Charlie", age: 35, status: "A" },
  { name: "David", age: 40, status: "B" },
  { name: "Eve", age: 45, status: "A" },
  { name: "Frank", age: 50, status: "B" },
  { name: "Grace", age: 55, status: "A" },
  { name: "Hank", age: 60, status: "B" },
  { name: "Irene", age: 65, status: "A" },
  { name: "Jack", age: 70, status: "B" }
])
```



OPERATIONS

1. AND

In MongoDB, you can use the comma `,` operator to perform an AND operation.

Syntax:

```
db.collection.find({ field1: value1, field2: value2 })
```

Example:

```
db.EX7.find({ age: 25, status: 'A' })
```

```
>_MONGOSH
> db.EX7.find({ age: 25, status: 'A' })
< {
  _id: ObjectId('65f6fc6287028d10f5c94385'),
  name: 'Alice',
  age: 25,
  status: 'A'
}
LAB>
```

2. OR

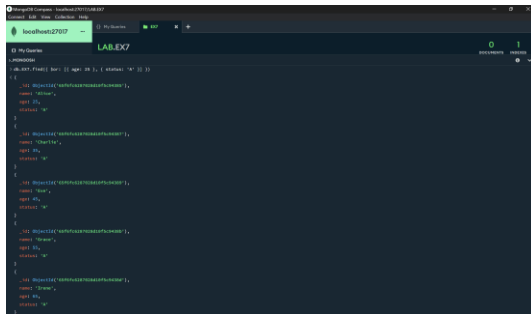
You can use the `\$or` operator to perform an OR operation.

Syntax:

```
db.collection.find({ $or: [{ field1: value1 }, { field2: value2 }] })
```

Example:

```
db.EX7.find({ $or: [{ age: 25 }, { status: 'A' }] })
```



```

> db.EX7.find({ $or: [{ age: 25 }, { status: 'A' }] })
[
  {
    "_id": "65f6fc6287028d10f5c94385",
    "name": "Alice",
    "age": 25,
    "status": "A"
  },
  {
    "_id": "65f6fc6287028d10f5c94386",
    "name": "Bob",
    "age": 30,
    "status": "B"
  },
  {
    "_id": "65f6fc6287028d10f5c94387",
    "name": "Charlie",
    "age": 25,
    "status": "A"
  }
]
```

3. LIMIT RECORDS

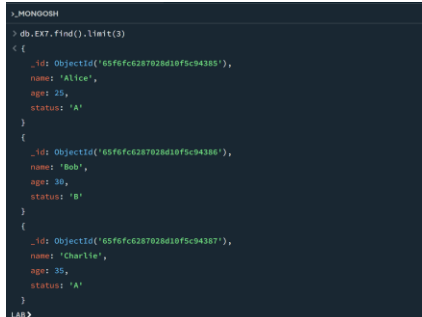
You can use the `limit()` method to limit the number of records.

Syntax:

```
db.collection.find().limit(number)
```

Example:

```
db.EX7.find().limit(3)
```



```

>_MONGOSH
> db.EX7.find().limit(3)
< [
  {
    _id: ObjectId('65f6fc6287028d10f5c94385'),
    name: 'Alice',
    age: 25,
    status: 'A'
  },
  {
    _id: ObjectId('65f6fc6287028d10f5c94386'),
    name: 'Bob',
    age: 30,
    status: 'B'
  },
  {
    _id: ObjectId('65f6fc6287028d10f5c94387'),
    name: 'Charlie',
    age: 25,
    status: 'A'
  }
]
LAB>
```

4. SORT RECORDS:

You can use the `sort()` method to sort the records.


Syntax:

```
db.collection.find().sort({ field1: 1 })
```

1 for ascending order, -1 for descending.

Example:

```
db.EX7.find().sort({ age: -1 })
```



5. INDEXING:

Indexing in MongoDB works by storing a small portion of the collection's data set in an easy to traverse form. The index stores the value of a specific field or set of fields, ordered by the value of the field.

Syntax:

```
db.collection.createIndex({ field1: 1 })
```

Example:

```
db.EX7.createIndex({ age: 1 })
```



6. ADVANCED INDEXING:

MongoDB also supports advanced indexing techniques like compound indexes, multikey indexes, text indexes, etc.

Syntax of creating a compound index:

```
db.collection.createIndex({ field1: 1, field2: -1 })
```

Example:

```
db.EX7.createIndex({ age: 1, status: -1 })
```




```

localhost:27017 ... {} My Queries EX7 x +
LAB.EX7
>_MONGOSH
> db.EX7.createIndex({ age: 1, status: -1 })
< age_1_status_-1
LAB >

```

7. AGGREGATION:

Aggregation operations process data records and return computed results. Aggregation operations group values from multiple documents together, and can perform a variety of operations on the grouped data to return a single result. MongoDB provides three ways to perform aggregation: the aggregation pipeline, the map-reduce function, and single purpose aggregation methods.

Syntax of an aggregation pipeline:

```

db.collection.aggregate([
  { $stage1: { ... } },
  { $stage2: { ... } },
  { $stage3: { ... } },
  { $stage4: { ... } },
  ...
], { allowDiskUse: true })

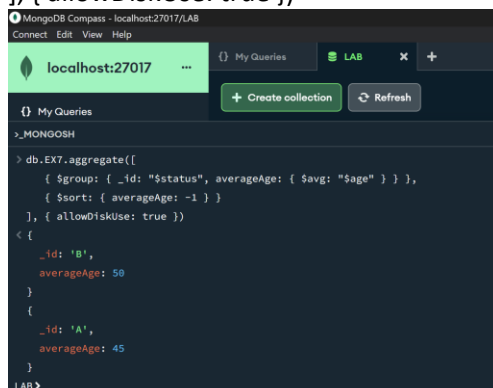
```

Example:

```

db.EX7.aggregate([
  { $group: { _id: "$status", averageAge: { $avg: "$age" } } },
  { $sort: { averageAge: -1 } }
], { allowDiskUse: true })

```



```

MongoDB Compass - localhost:27017/LAB
Connect Edit View Help
localhost:27017 ... {} My Queries LAB x +
+ Create collection Refresh
>_MONGOSH
> db.EX7.aggregate([
  { $group: { _id: "$status", averageAge: { $avg: "$age" } } },
  { $sort: { averageAge: -1 } }
], { allowDiskUse: true })
< {
  _id: 'B',
  averageAge: 50
}
{
  _id: 'A',
  averageAge: 45
}
LAB >

```

8. MAP REDUCE: (DEPRICATED/REMOVED FROM MongoDB)

Map-reduce is a data processing paradigm for condensing large volumes of data into useful aggregated results. MongoDB uses mapReduce command for map-reduce operations.

MongoDB has deprecated the mapReduce function in favor of using the aggregate function, which can perform the same operations more efficiently.

Syntax:

```
db.collection.mapReduce(  
  function() { emit(key, value); },  
  function(key, values) { return reduceFunction(values) },  
  {  
    out: outputCollection,  
    query: query, // Optional: query to select input documents  
    sort: sort, // Optional: sort the input documents  
    limit: limit // Optional: limit the number of input documents  
  }  
)
```

- ✓ **\$match:** Filters the documents to pass only documents that match the specified condition(s) to the next pipeline stage.
- ✓ **\$group:** Groups input documents by a specified identifier expression and applies the accumulator expression(s), if specified, to each group.
- ✓ **\$sort:** Reorders the document stream by a specified sort key. Only the order changes; the documents remain unmodified.
- ✓ **\$limit:** Passes the first "n" documents unmodified to the pipeline where "n" is the specified limit. For each input document, outputs either one document (unchanged) or zero documents.
- ✓ **\$out:** Writes the resulting documents of the aggregation pipeline to a collection

Exercise- 8

Column oriented databases study, queries, and practices

Column-oriented databases

Column-oriented databases, such as MongoDB, store data in columns rather than rows, which provides several advantages for analytical processing and reporting. MongoDB is a popular NoSQL document-oriented database that supports both row-oriented and column-oriented storage. It offers a flexible schema and supports various querying techniques, including document-based queries, aggregation pipelines, and indexing. Key aspects of column-oriented databases and MongoDB include:

1. **Comparison with row-oriented databases:** Column-oriented databases are optimized for analytical processing and reporting, while row-oriented databases are better suited for transactional processing. MongoDB supports both row-oriented and column-oriented storage, allowing it to handle a wide range of workloads.
2. **Performance advantages:** Column-oriented databases can achieve faster query performance by keeping data closer together and reducing seek time. MongoDB uses column-wise compression techniques to improve storage efficiency and query performance.
3. **Query examples:** MongoDB supports various querying techniques, including document-based queries, aggregation pipelines, and indexing. These features allow for efficient querying of data stored in columns.
4. **Best practices:** When working with column-oriented databases like MongoDB, it's essential to consider factors such as data modeling, indexing, and query optimization to ensure optimal performance. MongoDB provides tools and resources to help users optimize their databases for specific workloads.

In summary, column-oriented databases, such as MongoDB, offer several advantages for analytical processing and reporting, including faster query performance, efficient storage, and support for various querying techniques. By understanding the unique characteristics of column-oriented databases and best practices for working with them, users can effectively leverage these systems for their data management needs.

QUERIES

Column-oriented databases, such as MongoDB, offer several advantages for querying and processing data stored in columns. Here are some key aspects of querying in column-oriented databases:

1. **Document-based queries:** MongoDB supports document-based queries, which allow users to retrieve specific documents based on various criteria. These queries can be performed using MongoDB's query language, which supports a wide range of operators and filtering options.
2. **Aggregation pipelines:** MongoDB provides aggregation pipelines, which are a set of stages that can be applied to a collection of documents to perform various transformations and computations. These

pipelines can be used to perform complex queries, such as grouping, sorting, and filtering, on large volumes of data.

3. **Indexing:** Column-oriented databases, like MongoDB, support indexing to improve query performance. Indexes can be created on specific columns or fields, allowing for faster retrieval of data based on those columns.
4. **Query optimization:** MongoDB provides tools and resources to help users optimize their queries for specific workloads. This includes features like query profiling and explain plans, which allow users to analyze the performance of their queries and identify potential bottlenecks.
5. **Real-time analytics:** Column-oriented databases, such as MongoDB, are well-suited for real-time analytics and event-driven workloads. They can handle high volumes of data and provide fast query performance, making them suitable for applications that require real-time data processing and analysis.

In summary, column-oriented databases like MongoDB offer several advantages for querying and processing data stored in columns, including document-based queries, aggregation pipelines, indexing, query optimization, and support for real-time analytics. By understanding the unique features and capabilities of column-oriented databases, users can effectively leverage these systems for their data management and analysis needs.

Practises

For column-oriented databases like MongoDB, there are several practices and techniques that can be employed to optimize performance and improve query efficiency. Some of these practices include:

1. **Data modeling:** Choosing an appropriate data model is crucial for efficient data access and querying. MongoDB supports various data modeling approaches, such as document-oriented modeling and schema-less modeling. It is essential to consider factors like data size, query patterns, and performance requirements when designing the data model.
2. **Indexing:** Indexing is a critical aspect of MongoDB performance tuning. Indexes help speed up read operations by allowing the database to locate documents more efficiently. By creating appropriate indexes for frequently queried fields, you can significantly improve the performance of your queries. MongoDB supports various index types, including single-field indexes, compound indexes, and partial indexes.
3. **Query optimization:** Analyzing query performance and optimizing queries can help improve the efficiency of your database. MongoDB provides tools like the \$explain operator, which allows you to analyze query performance and identify potential bottlenecks. Additionally, optimizing query patterns, such as using projection to limit the returned fields and utilizing query operators and aggregation for efficient filtering and data manipulation, can help improve query performance.

4. **Performance tuning:** Tuning hardware and OS configuration, managing resources, and optimizing schema design can all contribute to better MongoDB performance. For example, selecting the right hardware, managing connection pools, and configuring read and write concerns can help balance performance and availability.
5. **Query execution:** Understanding the physical layout of column-oriented databases, such as how data is stored and accessed, can help optimize query performance. Techniques like tuple construction and join processing can be optimized to improve I/O and CPU utilization.

By following these practices and techniques, you can effectively optimize the performance of your column-oriented database, such as MongoDB, and achieve better query efficiency and faster data retrieval.

Exercise- 9

Create a database that stores road cars. Cars have a manufacturer, a type. Each car has a maximum performance and a maximum torque value. Do the following: Test Cassandras replication schema and consistency models..

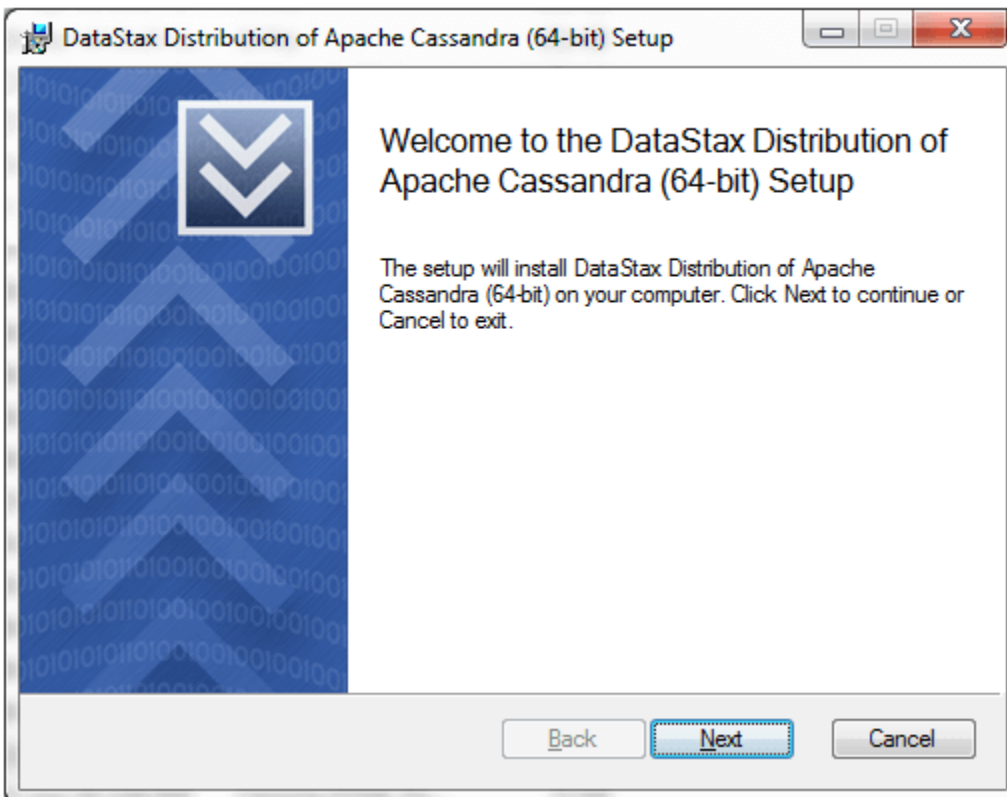
Installation : Download the Cassandra software from the link below

<https://datastax-community-edition.software.informer.com/download/>

Download and Install Cassandra

Follow the steps given below:

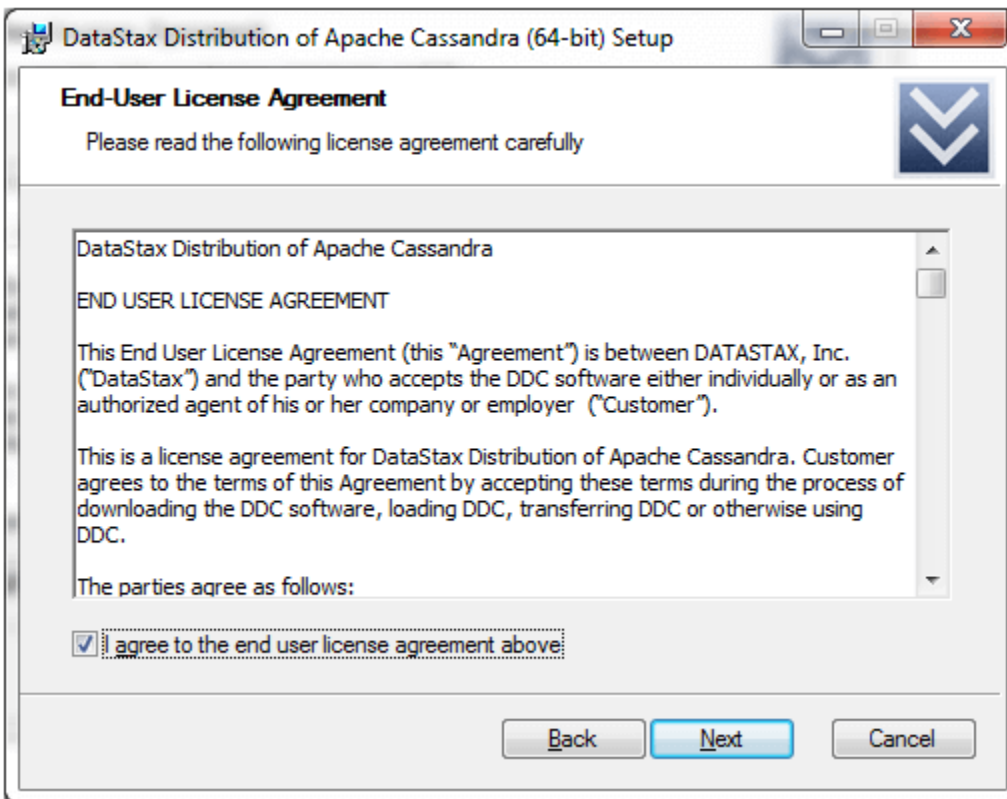
- Run the datastax community edition setup. After running the setup, you will see the following page will be displayed. It is a screenshot of 64 bit version.



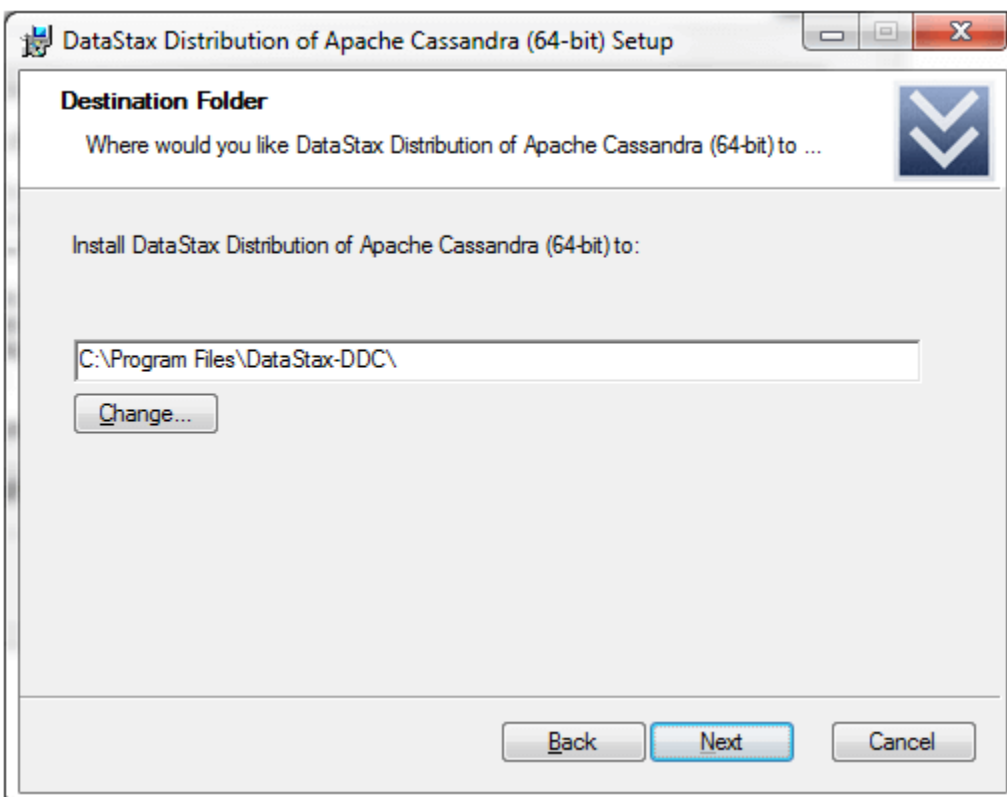
- Click on the next button and you will get the following page:

Register No.....

Date:.....



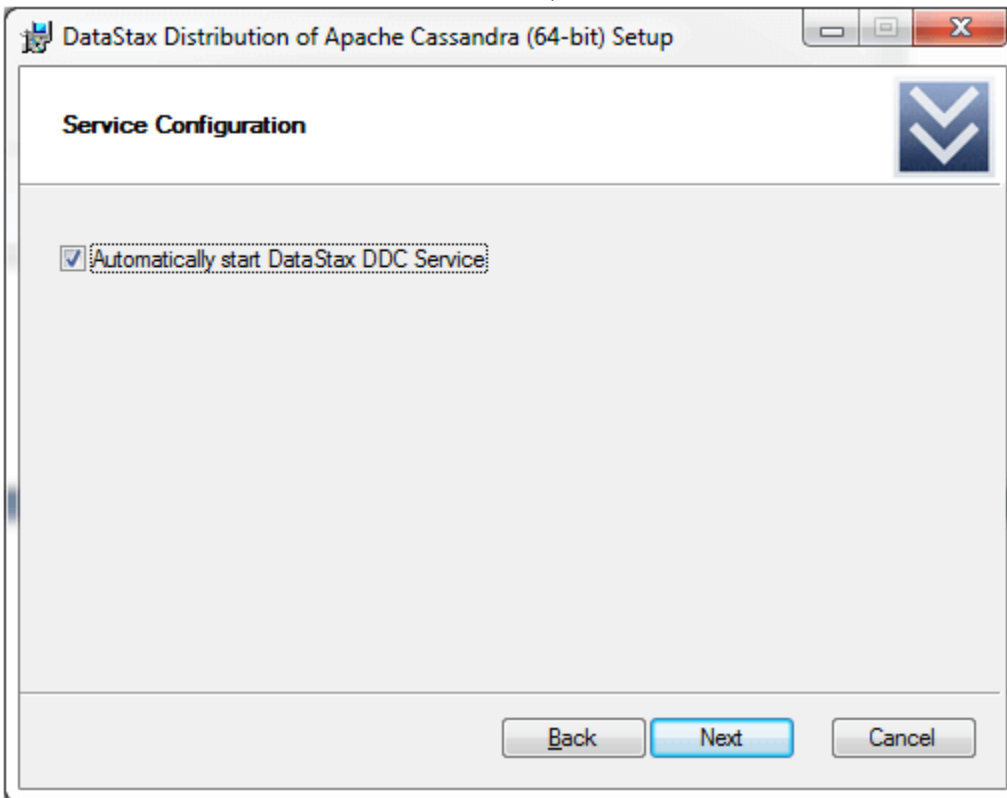
- Press the next button and you will get the following page specifying the location of the installation.

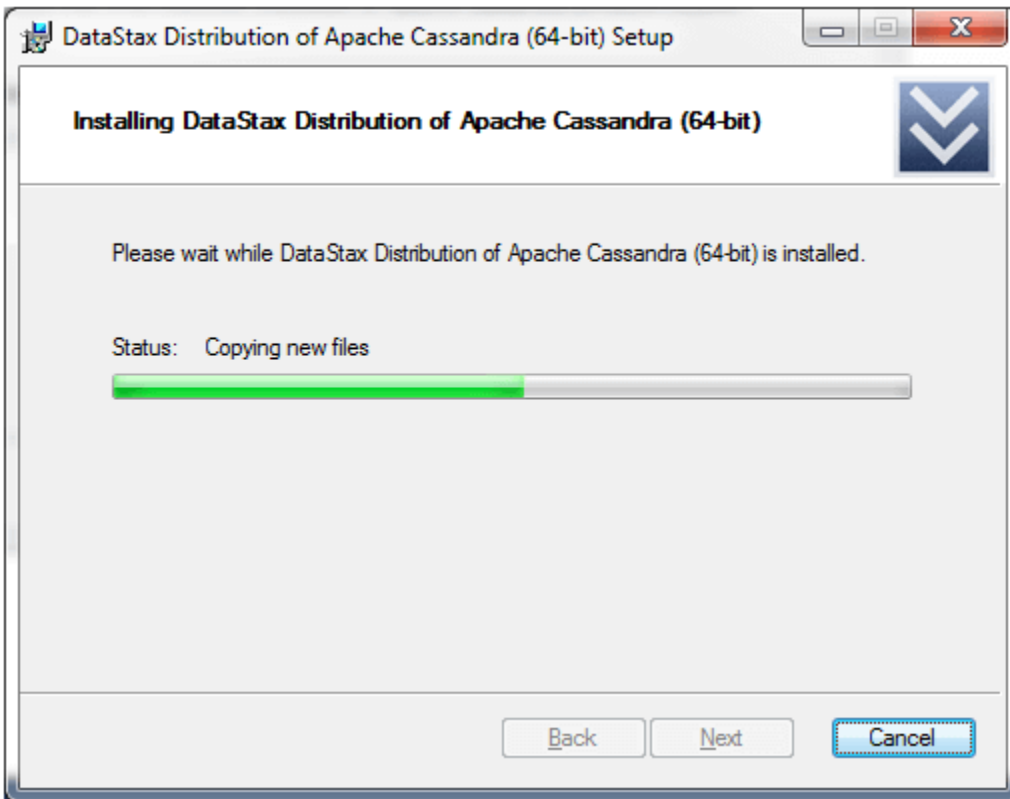


Register No.

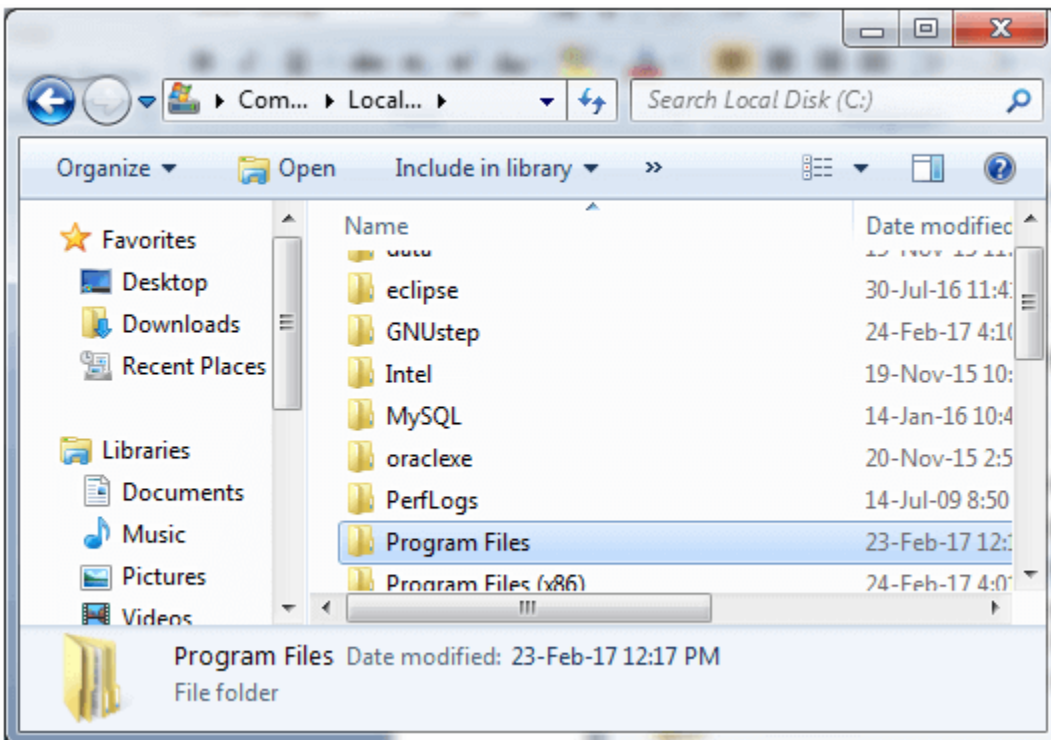
Date:

- Press the next button and a page will appear asking about whether you automatically start Data Stax DDC service. Click on the radio button and proceed next.

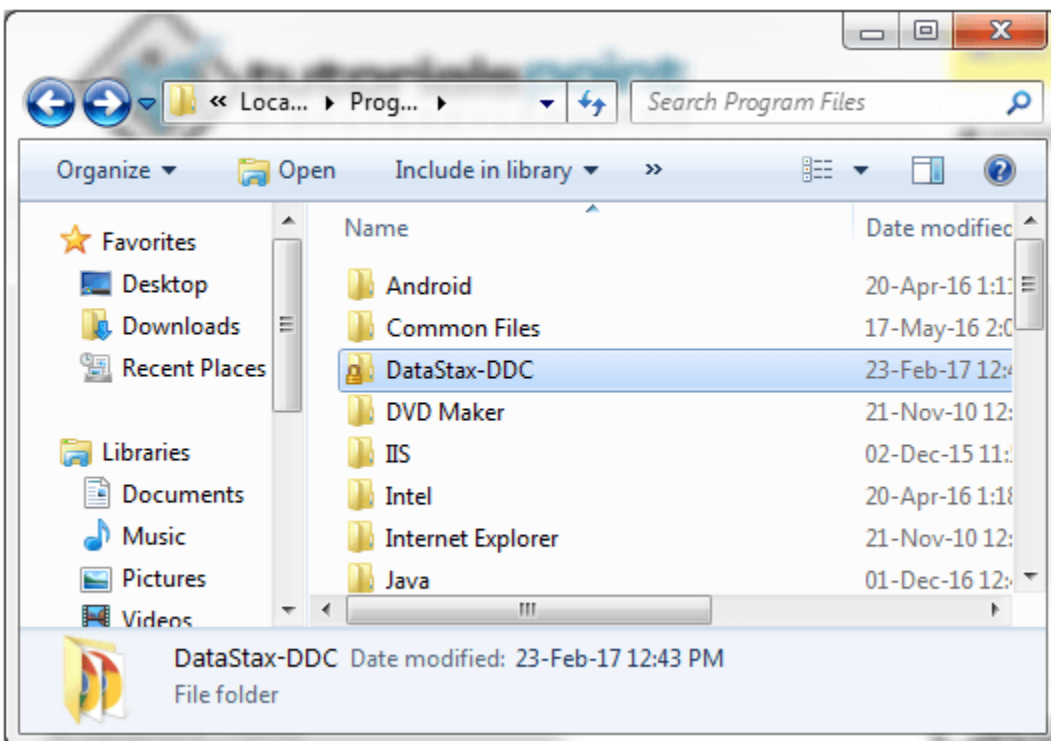




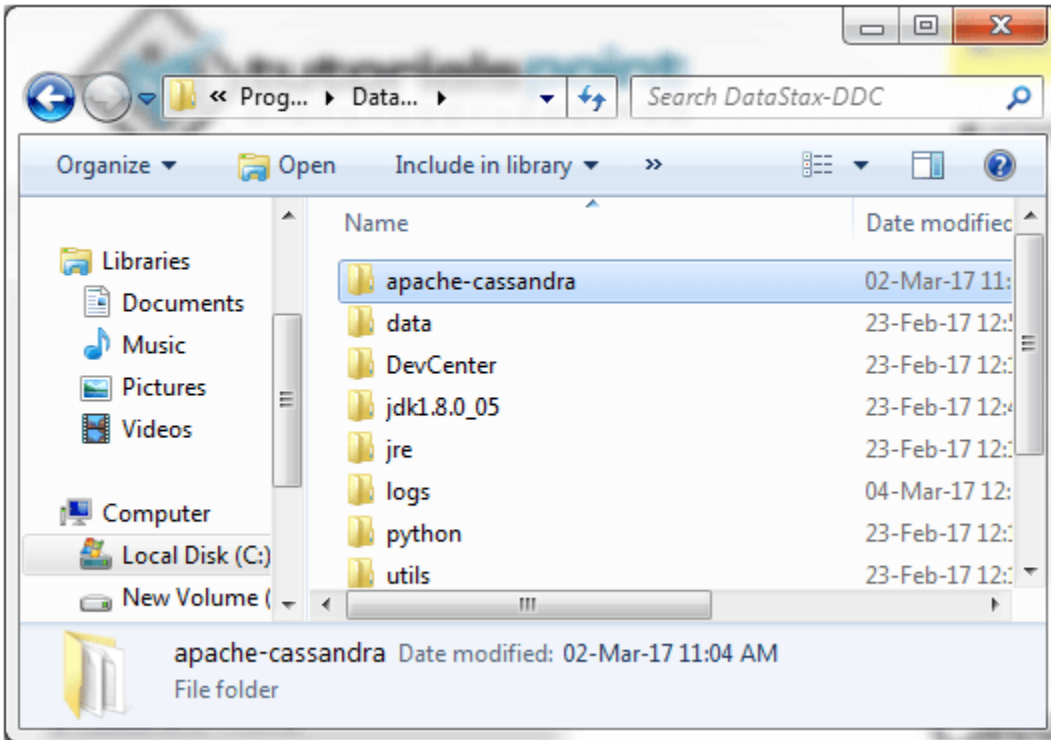
- Installation is started now. After completing the installation, go to program files where Data Stax is installed.



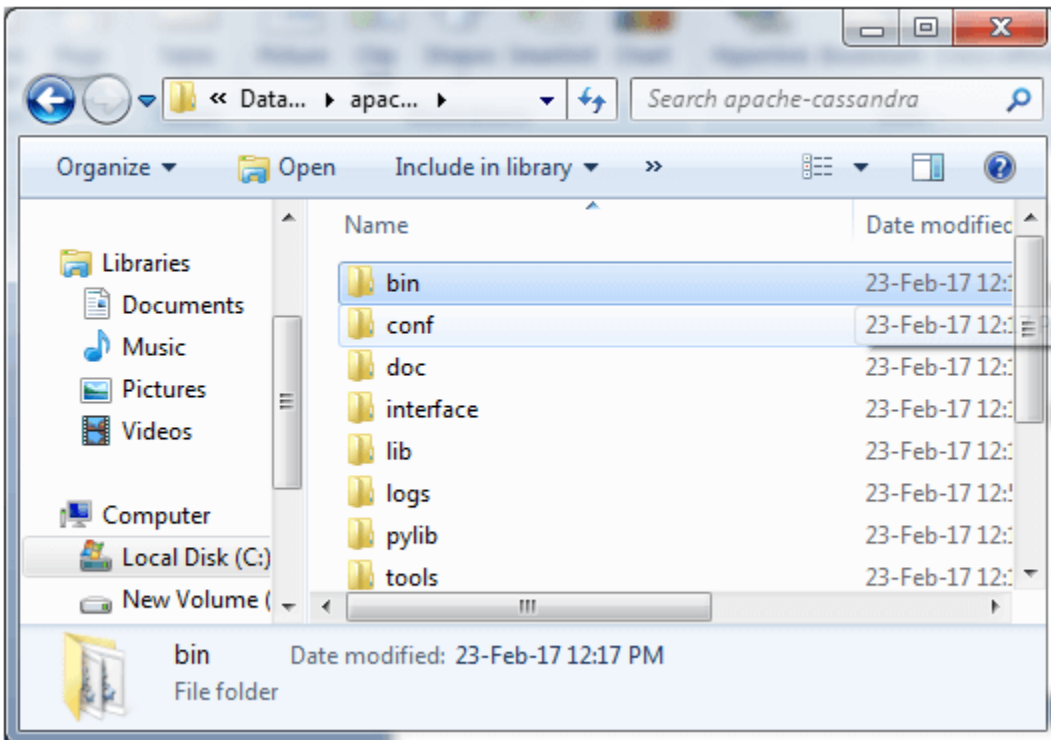
- Open Program Files then you see the following page:



-
- Open DataStax-DDC then you see Apache Cassandra:




- Open Apache Cassandra and you see bin:



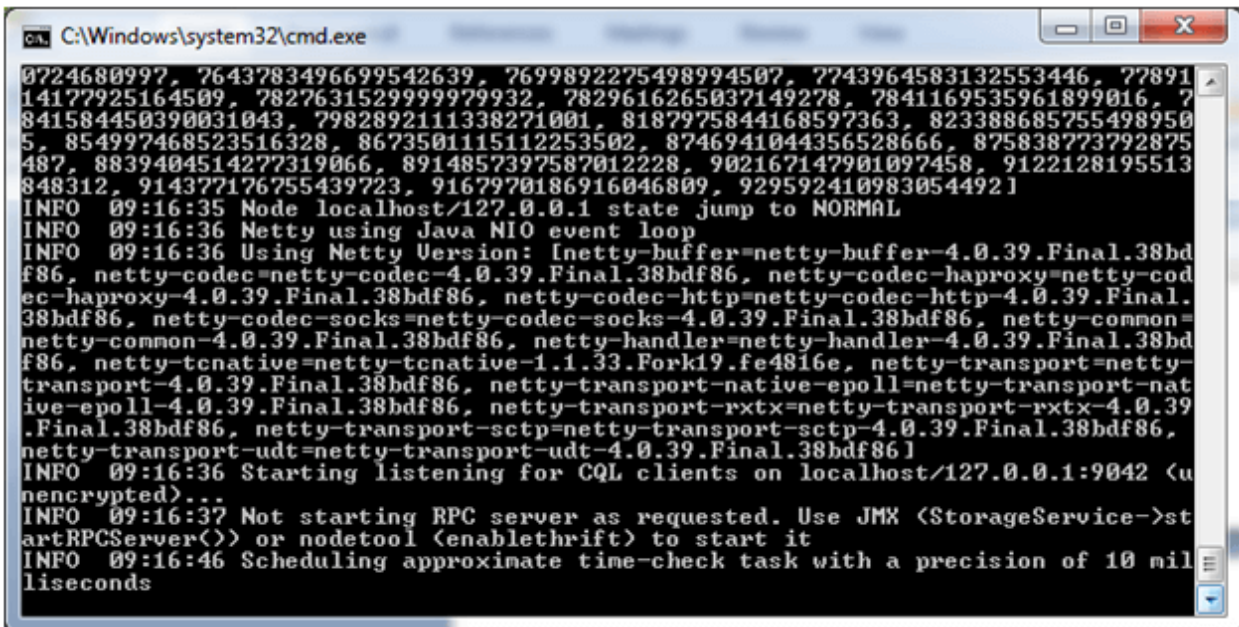
- Open bin and you will see Cassandra Windows batch File:

Register No.....

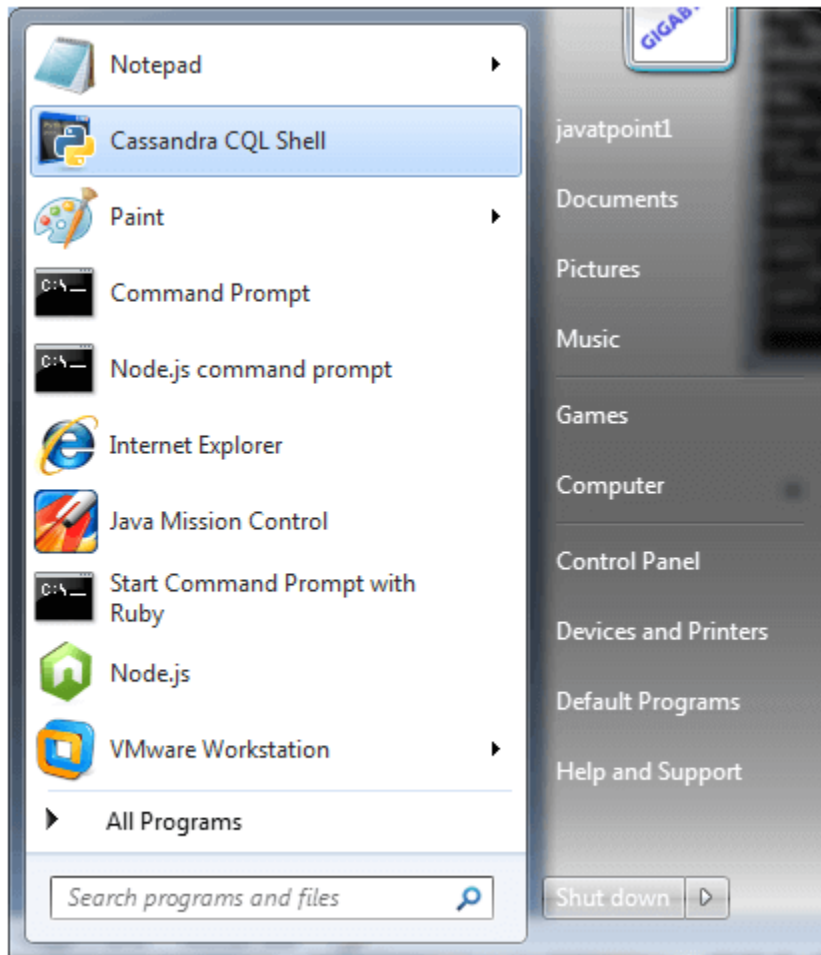
Date:.....




- Run this file. It will start the Cassandra server and you will see the following page:



- Server is started now go to windows start programs, search Cassandra Shell



- Run the Cassandra Shell. After running Cassandra shell, you will see the following command line:



Description

Step 1: Define Keyspace

In Cassandra, a keyspace is like a schema in relational databases. It defines the replication strategy and other options for data distribution.

```
CREATE KEYSPACE IF NOT EXISTS car_keyspace  
WITH replication = {'class': 'SimpleStrategy', 'replication_factor': 3};
```

In this example, we are using SimpleStrategy with a replication factor of 3 for simplicity. In a production environment, you would adjust the replication strategy based on your needs.

Step 2: Create Table

Now, let's create a table to store information about road cars.

```
CREATE TABLE IF NOT EXISTS car_keyspace.cars (  
    car_id UUID PRIMARY KEY,  
    manufacturer TEXT,  
    type TEXT,  
    max_performance DOUBLE,  
    max_torque DOUBLE  
);
```

This table has columns for car_id (primary key), manufacturer, type, max_performance, and max_torque.

Step 3: Insert Data

You can insert data into the table using INSERT statements.

```
INSERT INTO car_keyspace.cars (car_id, manufacturer, type, max_performance, max_torque)
VALUES (uuid(), 'Toyota', 'SUV', 300, 400);
```

```
INSERT INTO car_keyspace.cars (car_id, manufacturer, type, max_performance, max_torque)
VALUES (uuid(), 'BMW', 'Sedan', 350, 450);
```

```
INSERT INTO car_keyspace.cars (car_id, manufacturer, type, max_performance, max_torque)
VALUES (uuid(), 'Ford', 'Truck', 250, 350);
```

Step 4: Query Data

You can retrieve data from the table using SELECT statements.

```
SELECT * FROM car_keyspace.cars
```

Step 5: Replication and Consistency

Cassandra provides various replication strategies and consistency levels. By default, Cassandra uses eventual consistency, but you can configure consistency levels on a per-query basis or globally.

Replication can be tested by setting up a Cassandra cluster with multiple nodes and observing how data is replicated across nodes.

Consistency can be tested by performing read and write operations with different consistency levels and observing the results.

For example:

```
INSERT INTO car_keyspace.cars (car_id, manufacturer, type, max_performance, max_torque)
VALUES (uuid(), 'Mercedes', 'Coupe', 400, 500) USING CONSISTENCY QUORUM;

SELECT * FROM car_keyspace.cars USING CONSISTENCY QUORUM;
```

In this example, the INSERT statement specifies a consistency level of QUORUM, ensuring that the write operation is acknowledged by a majority of replicas. The subsequent SELECT statement also uses QUORUM consistency level, ensuring that the data retrieved is consistent across replicas.

By experimenting with different consistency levels and observing the results, you can test Cassandra's consistency models effectively.

Program

Step 1

```
CREATE KEYSPACE IF NOT EXISTS car_keyspace
```

```
... WITH replication = {'class': 'SimpleStrategy', 'replication_factor': 3};
```

```
sh> CREATE TABLE IF NOT EXISTS car_keyspace.cars (
```

```
...   car_id UUID PRIMARY KEY,
```

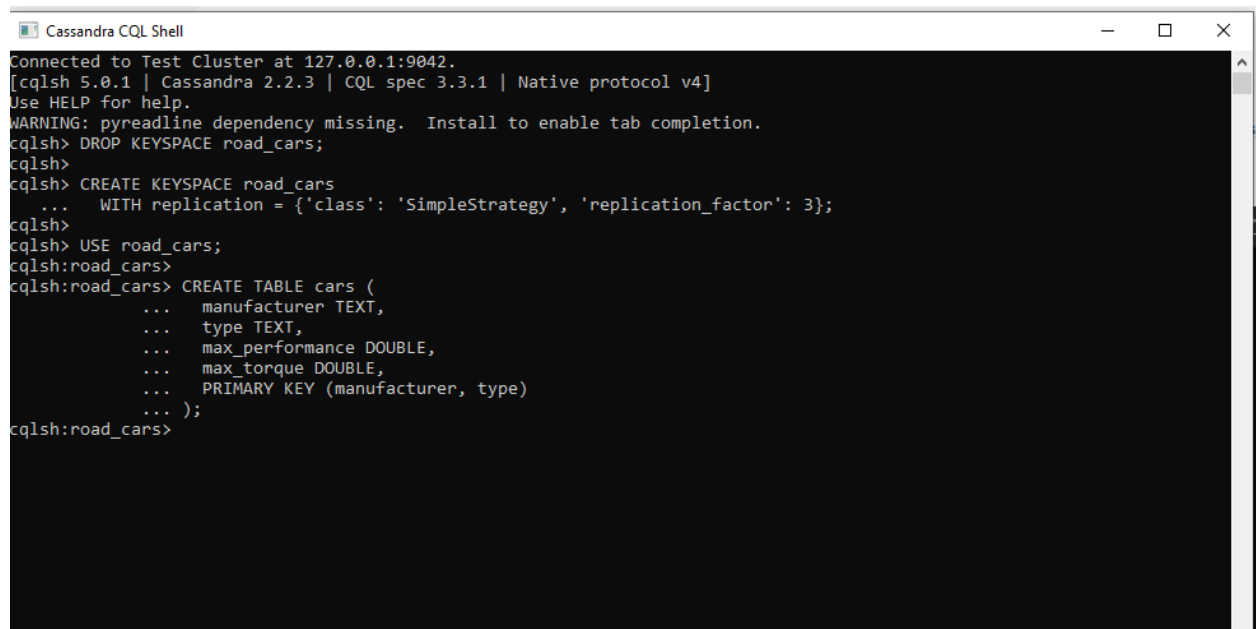
```
...   manufacturer TEXT,
```

```
...   type TEXT,
```

```
...   max_performance DOUBLE,
```

```
...   max_torque DOUBLE
```

```
... );
```

A screenshot of a terminal window titled "Cassandra CQL Shell". The terminal shows the following commands and output:

```
Connected to Test Cluster at 127.0.0.1:9042.
[cqlsh 5.0.1 | Cassandra 2.2.3 | CQL spec 3.3.1 | Native protocol v4]
Use HELP for help.
WARNING: pyreadline dependency missing.  Install to enable tab completion.
cqlsh> DROP KEYSPACE road_cars;
cqlsh>
cqlsh> CREATE KEYSPACE road_cars
...   WITH replication = {'class': 'SimpleStrategy', 'replication_factor': 3};
cqlsh>
cqlsh> USE road_cars;
cqlsh:road_cars>
cqlsh:road_cars> CREATE TABLE cars (
...   manufacturer TEXT,
...   type TEXT,
...   max_performance DOUBLE,
...   max_torque DOUBLE,
...   PRIMARY KEY (manufacturer, type)
... );
cqlsh:road_cars>
```

```
sh> INSERT INTO car_keyspace.cars (car_id, manufacturer, type, max_performance, max_torque)
... VALUES (uuid(), 'Toyota', 'SUV', 300, 400);
```

```
sh>
```

```
sh> INSERT INTO car_keyspace.cars (car_id, manufacturer, type, max_performance, max_torque)
```



```
... VALUES (uuid(), 'BMW', 'Sedan', 350, 450);
```

```
sh>
```

```
sh> INSERT INTO car_keyspace.cars (car_id, manufacturer, type, max_performance, max_torque)
```

```
... VALUES (uuid(), 'Ford', 'Truck', 250, 350);
```

```
cqlsh:road_cars> INSERT INTO cars (manufacturer, type, max_performance, max_torque)
... VALUES ('Toyota', 'Sedan', 200, 180);
cqlsh:road_cars>
cqlsh:road_cars> INSERT INTO cars (manufacturer, type, max_performance, max_torque)
... VALUES ('Honda', 'SUV', 220, 200);
cqlsh:road_cars>
cqlsh:road_cars> INSERT INTO cars (manufacturer, type, max_performance, max_torque)
... VALUES ('Ford', 'Truck', 250, 220);
```

```
sh> SELECT * FROM car_keyspace.cars;
```

car_id	manufacturer	max_performance	max_torque	type
29c773a0-dc1c-4d54-88b5-4b93d4a98d10	Toyota	300	400	SUV
1ee85e4f-2f75-4afb-9724-e0d9135ed846	Ford	250	350	Truck
b78950ae-d043-40b6-a6ac-166293910070	BMW	350	450	Sedan

```
car_id | manufacturer | max_performance | max_torque | type
-----+-----+-----+-----+-----
29c773a0-dc1c-4d54-88b5-4b93d4a98d10 | Toyota | 300 | 400 | SUV
1ee85e4f-2f75-4afb-9724-e0d9135ed846 | Ford | 250 | 350 | Truck
b78950ae-d043-40b6-a6ac-166293910070 | BMW | 350 | 450 | Sedan
(3 rows)
cqlsh:road_cars>
```

```
(3 rows)
```

Exercise– 10

Master Data Management using Neo4j Manage your master data more effectively the world of master data is changing. Data architects and application developers are swapping their relational databases with graph databases to store their master data. This switch enables them to use a data store optimized to discover new insights in existing data, provide a 360-degree view of master data and answer questions about data relationships in real time write an essay with examples and explanation with output

Title: Unleashing the Power of Neo4j for Master Data Management

In the dynamic landscape of data management, the traditional paradigm of using relational databases for master data management is undergoing a paradigm shift. Data architects and application developers are increasingly turning to graph databases, notably Neo4j, to store and manage their master data. This transformation heralds a new era of master data management, one characterized by enhanced agility, deeper insights, and real-time relationship exploration. This essay explores the evolution of master data management with Neo4j, elucidating its benefits, and providing examples to illustrate its transformative capabilities.

Understanding Neo4j for Master Data Management:

Neo4j is a leading graph database that excels in representing and querying highly connected data. Unlike relational databases, which rely on tabular structures, Neo4j employs a graph-based model consisting of nodes, relationships, and properties. This model is particularly well-suited for master data management, as it enables the representation of complex relationships between entities in a natural and intuitive manner.

Benefits of Neo4j for Master Data Management:

1. Discovering New Insights: By leveraging the graph-based nature of Neo4j, organizations can uncover hidden relationships and patterns within their master data. For example, in a customer master data scenario, Neo4j can reveal previously unrecognized connections between customers based on shared attributes, interactions, or transactions.

2. 360-Degree View of Master Data: Neo4j facilitates the creation of a holistic view of master data by capturing not only the entities themselves but also the relationships between them. This enables organizations to gain a comprehensive understanding of their data landscape, facilitating informed decision-making and strategic planning.

3. Real-Time Relationship Exploration: With Neo4j, organizations can explore data relationships in real-time, enabling dynamic analysis and visualization of interconnected data. For instance, in a product master data scenario, Neo4j can provide immediate insights into product hierarchies, dependencies, and associations, empowering organizations to adapt quickly to changing market dynamics.

Example: Customer Master Data Management with Neo4j:

Consider a scenario where a retail company seeks to manage its customer master data using Neo4j. The company stores information about customers, their purchases, and interactions across various touchpoints. Let's explore how Neo4j can facilitate this process:

```
``cypher
// Create nodes for customers and products
CREATE (:Customer {id: 1, name: 'Alice'}),
      (:Customer {id: 2, name: 'Bob'}),
      (:Product {id: 'P1', name: 'Laptop'}),
      (:Product {id: 'P2', name: 'Smartphone'})

// Create relationships between customers and products
MATCH (c:Customer {name: 'Alice'}), (p:Product {id: 'P1'})
CREATE (c)-[:PURCHASED]->(p)

MATCH (c:Customer {name: 'Bob'}), (p:Product {id: 'P2'})
CREATE (c)-[:PURCHASED]->(p)
...
```

In this example, we create nodes representing customers and products, and establish relationships between customers and the products they have purchased. This graph-based representation enables the

company to easily query and analyze customer purchase patterns, identify cross-selling opportunities, and personalize marketing campaigns based on customer preferences and behavior.

Conclusion:

In conclusion, Neo4j represents a paradigm shift in master data management, offering organizations a powerful platform to store, manage, and analyze their master data. By leveraging the inherent capabilities of graph databases, Neo4j enables organizations to gain deeper insights, create holistic views of their data, and explore data relationships in real-time. Through examples like customer master data management, it becomes evident that Neo4j's graph-based approach revolutionizes traditional master data management practices, empowering organizations to unlock new possibilities and drive innovation in the digital age.

Exercise– 11

Shopping Mall case study using Cassandra, where we have many customers ordering items from thermall and, we have suppliers who deliver them their ordered items

AIM: Shopping Mall case study using Cassandra, where we have many customers ordering items from thermall and, we have suppliers who deliver them their ordered items.

Description: In this lab experiment, we explore database management with Apache Cassandra. They begin by creating a keyspace to organize data, then establish tables for customers, suppliers, products, orders, and order details. Through CQL commands, they insert sample data representing customers, suppliers, and products. The experiment delves into querying techniques, enabling students to retrieve specific customer orders and analyze relationships between entities. Through hands-on practice, we gain proficiency in database schema design, data insertion, and retrieval, while comprehending the fundamentals of distributed database systems exemplified by Cassandra's schema flexibility and **eventual consistency model**.

Procedure:

Step 1: Once you run this command, it will remove the keyspace shopping_mall along with all the tables and data inside it. Make sure you're certain about dropping the keyspace as it will delete all data associated with it.

```
Connected to Test Cluster at 127.0.0.1:9042.  
[cqlsh 5.0.1 | Cassandra 2.2.3 | CQL spec 3.3.1 | Native protocol v4]  
Use HELP for help.  
WARNING: pyreadline dependency missing. Install to enable tab completion.  
cqlsh> -- Drop the keyspace if it exists  
cqlsh> DROP KEYSPACE IF EXISTS shopping_mall;  
cqlsh>
```

Step 2: Create a keyspace, define a table for customer data, and insert 5 customer records into a Cassandra database using CQL (Cassandra Query Language) in the Cassandra CQL shell:

```

cqlsh> -- Create a keyspace
cqlsh> CREATE KEYSPACE IF NOT EXISTS shopping_mall WITH replication = {'class': 'SimpleStrategy', 'replication_factor': 1};
cqlsh>
cqlsh> -- Use the keyspace
cqlsh> USE shopping_mall;
cqlsh:shopping_mall>
cqlsh:shopping_mall> -- Create a table to store customer data
cqlsh:shopping_mall> CREATE TABLE IF NOT EXISTS customers (
...     customer_id UUID PRIMARY KEY,
...     name TEXT,
...     email TEXT,
...     phone TEXT,
...     address TEXT
... );
cqlsh:shopping_mall>
cqlsh:shopping_mall> -- Insert 5 customers data
cqlsh:shopping_mall> INSERT INTO customers (customer_id, name, email, phone, address) VALUES (uuid(), 'John Doe', 'john.doe@example.com', '123-456-7890', '123 Main St');
cqlsh:shopping_mall> INSERT INTO customers (customer_id, name, email, phone, address) VALUES (uuid(), 'Jane Smith', 'jane.smith@example.com', '456-789-0123', '456 Elm St');
cqlsh:shopping_mall> INSERT INTO customers (customer_id, name, email, phone, address) VALUES (uuid(), 'Alice Johnson', 'alice.johnson@example.com', '789-012-3456', '789 Oak St');
cqlsh:shopping_mall> INSERT INTO customers (customer_id, name, email, phone, address) VALUES (uuid(), 'Bob Brown', 'bob.brown@example.com', '012-345-6789', '012 Pine St');
cqlsh:shopping_mall> INSERT INTO customers (customer_id, name, email, phone, address) VALUES (uuid(), 'Emily Davis', 'emily.davis@example.com', '345-678-9012', '345 Cedar St');
cqlsh:shopping_mall>
cqlsh:shopping_mall> -- Query the data to verify insertion
cqlsh:shopping_mall> SELECT * FROM customers;

```

customer_id	address	email	name	phone
183525a4-f944-4561-831a-9c5d2c110382	456 Elm St	jane.smith@example.com	Jane Smith	456-789-0123
260e336c-dc47-42db-85f8-99c08e5c3a57	345 Cedar St	emily.davis@example.com	Emily Davis	345-678-9012
22f87bcd-861f-47c5-af64-8c18690d5665	123 Main St	john.doe@example.com	John Doe	123-456-7890
5a31d548-1fe3-4bca-a795-47d07eab94d2	012 Pine St	bob.brown@example.com	Bob Brown	012-345-6789
0be5f4a1-7ec3-4d57-914f-39e3b7fab853	789 Oak St	alice.johnson@example.com	Alice Johnson	789-012-3456

(5 rows)

This code will create a keyspace named shopping_mall, a table named customers to store customer data, and insert 5 sample customer records into the table. Finally, it selects and displays all the customer data to verify that the insertion was successful.

Step 3: This code will create a table named suppliers to store supplier data, and insert 5 sample supplier records into the table. Finally, it selects and displays all the supplier data to verify that the insertion was successful.

```

cqlsh:shopping_mall> -- Create a keyspace
cqlsh:shopping_mall> CREATE KEYSPACE IF NOT EXISTS shopping_mall WITH replication = {'class': 'SimpleStrategy', 'replication_factor': 1};
cqlsh:shopping_mall>
cqlsh:shopping_mall> -- Use the keyspace
cqlsh:shopping_mall> USE shopping_mall;
cqlsh:shopping_mall>
cqlsh:shopping_mall> -- Create a table to store supplier data
cqlsh:shopping_mall> CREATE TABLE IF NOT EXISTS suppliers (
...     supplier_id UUID PRIMARY KEY,
...     name TEXT,
...     email TEXT,
...     phone TEXT,
...     address TEXT
... );
cqlsh:shopping_mall>
cqlsh:shopping_mall> -- Insert 5 supplier records
cqlsh:shopping_mall> INSERT INTO suppliers (supplier_id, name, email, phone, address) VALUES (uuid(), 'Supplier 1', 'supplier1@example.com', '123-456-7890', '123 Supplier St');
cqlsh:shopping_mall> INSERT INTO suppliers (supplier_id, name, email, phone, address) VALUES (uuid(), 'Supplier 2', 'supplier2@example.com', '456-789-0123', '456 Supplier St');
cqlsh:shopping_mall> INSERT INTO suppliers (supplier_id, name, email, phone, address) VALUES (uuid(), 'Supplier 3', 'supplier3@example.com', '789-012-3456', '789 Supplier St');
cqlsh:shopping_mall> INSERT INTO suppliers (supplier_id, name, email, phone, address) VALUES (uuid(), 'Supplier 4', 'supplier4@example.com', '012-345-6789', '012 Supplier St');
cqlsh:shopping_mall> INSERT INTO suppliers (supplier_id, name, email, phone, address) VALUES (uuid(), 'Supplier 5', 'supplier5@example.com', '345-678-9012', '345 Supplier St');
cqlsh:shopping_mall>
cqlsh:shopping_mall> -- Query the data to verify insertion
cqlsh:shopping_mall> SELECT * FROM suppliers;

```

supplier_id	address	email	name	phone
d5aa00d9-5642-4eb0-b6c9-59190a43f22b	789 Supplier St	supplier3@example.com	Supplier 3	789-012-3456
e4a6cd3f-09f4-4577-8822-668f61aa74af	012 Supplier St	supplier4@example.com	Supplier 4	012-345-6789
3a645acc-a98c-416f-a4af-e58c450150ea	456 Supplier St	supplier2@example.com	Supplier 2	456-789-0123
5b4beab8-712b-4276-a934-4f4d4c546255	123 Supplier St	supplier1@example.com	Supplier 1	123-456-7890
27c48f00-c5a4-49f1-882f-a66b0339791a	345 Supplier St	supplier5@example.com	Supplier 5	345-678-9012

(5 rows)

Step 4:

Below is the CQL code to create a table for products and insert 5 different product records into the Cassandra database. This code will create a table named products to store product data, and insert 5 different product records into the table. Finally, it selects and displays all the product data to verify that the insertion was successful.

```
cqlsh:shopping_mall> -- Create a keyspace
cqlsh:shopping_mall> CREATE KEYSPACE IF NOT EXISTS shopping_mall WITH replication = {'class': 'SimpleStrategy', 'replication_factor': 1};
cqlsh:shopping_mall>
cqlsh:shopping_mall> -- Use the keyspace
cqlsh:shopping_mall> USE shopping_mall;
cqlsh:shopping_mall>
cqlsh:shopping_mall> -- Create a table to store product data
cqlsh:shopping_mall> CREATE TABLE IF NOT EXISTS products (
...     product_id UUID PRIMARY KEY,
...     name TEXT,
...     description TEXT,
...     price DECIMAL,
...     quantity INT
... );
cqlsh:shopping_mall>
cqlsh:shopping_mall> -- Insert 5 different product records
cqlsh:shopping_mall> INSERT INTO products (product_id, name, description, price, quantity) VALUES (uuid(), 'Product 1', 'Description of Product 1', 10.99, 100);
cqlsh:shopping_mall> INSERT INTO products (product_id, name, description, price, quantity) VALUES (uuid(), 'Product 2', 'Description of Product 2', 19.99, 50);
cqlsh:shopping_mall> INSERT INTO products (product_id, name, description, price, quantity) VALUES (uuid(), 'Product 3', 'Description of Product 3', 5.99, 200);
cqlsh:shopping_mall> INSERT INTO products (product_id, name, description, price, quantity) VALUES (uuid(), 'Product 4', 'Description of Product 4', 29.99, 75);
cqlsh:shopping_mall> INSERT INTO products (product_id, name, description, price, quantity) VALUES (uuid(), 'Product 5', 'Description of Product 5', 15.99, 150);
cqlsh:shopping_mall>
cqlsh:shopping_mall> -- Query the data to verify insertion
cqlsh:shopping_mall> SELECT * FROM products;
```

product_id	description	name	price	quantity
8c321e4d-150c-4156-acc8-e2bc5e432da9	Description of Product 2	Product 2	19.99	50
529d8669-5b2d-4931-b6f2-76662d5d890c	Description of Product 1	Product 1	10.99	100
d2a1e32d-c7b2-46eb-ad31-ad7903f7fd41	Description of Product 5	Product 5	15.99	150
324e6baa-75f5-4b34-a09d-f03de75433c5	Description of Product 3	Product 3	5.99	200
111723b3-2d0b-41b7-af20-e16f8ebf1079	Description of Product 4	Product 4	29.99	75

(5 rows)

Step 5:

A scenario where 5 customers are ordering products from suppliers, we'll need to first define tables for orders and order details. Below is the CQL code to set up the scenario. In this scenario, we create two tables: orders to store order information and order details to store details of each product ordered in an order. We then insert 5 orders into the orders table, each representing a different customer's order. Additionally, we insert order details for each order into the order details table.

```

cqlsh:shopping_mall> -- Create a keyspace
cqlsh:shopping_mall> CREATE KEYSPACE IF NOT EXISTS shopping_mall WITH replication = {'class': 'SimpleStrategy', 'replication_factor': 1};
cqlsh:shopping_mall>
cqlsh:shopping_mall> -- Use the keyspace
cqlsh:shopping_mall> USE shopping_mall;
cqlsh:shopping_mall>
cqlsh:shopping_mall> -- Create a table to store orders
cqlsh:shopping_mall> CREATE TABLE IF NOT EXISTS orders (
...     order_id UUID PRIMARY KEY,
...     customer_id UUID,
...     order_date TIMESTAMP,
...     total_price DECIMAL
... );
cqlsh:shopping_mall>
cqlsh:shopping_mall> -- Create a table to store order details
cqlsh:shopping_mall> CREATE TABLE IF NOT EXISTS order_details (
...     order_id UUID,
...     product_id UUID,
...     quantity INT,
...     price DECIMAL,
...     PRIMARY KEY (order_id, product_id)
... );
cqlsh:shopping_mall>
cqlsh:shopping_mall> -- Insert 5 different customer records
cqlsh:shopping_mall> INSERT INTO orders (order_id, customer_id, order_date, total_price) VALUES (uuid(), uuid(), toTimestamp(now()), 0);
cqlsh:shopping_mall> INSERT INTO orders (order_id, customer_id, order_date, total_price) VALUES (uuid(), uuid(), toTimestamp(now()), 0);
cqlsh:shopping_mall> INSERT INTO orders (order_id, customer_id, order_date, total_price) VALUES (uuid(), uuid(), toTimestamp(now()), 0);
cqlsh:shopping_mall> INSERT INTO orders (order_id, customer_id, order_date, total_price) VALUES (uuid(), uuid(), toTimestamp(now()), 0);
cqlsh:shopping_mall> INSERT INTO orders (order_id, customer_id, order_date, total_price) VALUES (uuid(), uuid(), toTimestamp(now()), 0);
cqlsh:shopping_mall>
cqlsh:shopping_mall> -- Insert order details for each order
cqlsh:shopping_mall> INSERT INTO order_details (order_id, product_id, quantity, price) VALUES (uuid(), uuid(), 2, 10.99);
cqlsh:shopping_mall> INSERT INTO order_details (order_id, product_id, quantity, price) VALUES (uuid(), uuid(), 3, 19.99);
cqlsh:shopping_mall> INSERT INTO order_details (order_id, product_id, quantity, price) VALUES (uuid(), uuid(), 1, 5.99);
cqlsh:shopping_mall> INSERT INTO order_details (order_id, product_id, quantity, price) VALUES (uuid(), uuid(), 4, 29.99);
cqlsh:shopping_mall> INSERT INTO order_details (order_id, product_id, quantity, price) VALUES (uuid(), uuid(), 2, 15.99);
cqlsh:shopping_mall>

```

Sample Queries

1. Query to Retrieve All Customers and its output

```

cqlsh:shopping_mall> SELECT * FROM customers;

```

customer_id	address	email	name	phone
d137a07b-ba7b-4796-ba40-693be86f6351	789 Oak St	alice.johnson@example.com	Alice Johnson	789-012-3456
fb7fa777-ad22-48c1-932a-9195ed01de60	123 Main St	john.doe@example.com	John Doe	123-456-7890
5c2bdeb8-acd3-462c-888b-d9cc9b719b31	456 Elm St	jane.smith@example.com	Jane Smith	456-789-0123
85461be9-71bf-4efe-a8f8-794c6e7ca034	345 Cedar St	emily.davis@example.com	Emily Davis	345-678-9012
c1f54bdc-6304-4bca-b0f9-55b17125f8ae	012 Pine St	bob.brown@example.com	Bob Brown	012-345-6789

2. Query to Retrieve All Products and its output


```
cqlsh:shopping_mall> SELECT * FROM products;
```

product_id	description	name	price	quantity
8c321e4d-150c-4156-acc8-e2bc5e432da9	Description of Product 2	Product 2	19.99	50
529d8669-5b2d-4931-b6f2-76662d5d890c	Description of Product 1	Product 1	10.99	100
d2a1e32d-c7b2-46eb-ad31-ad7903f7fd41	Description of Product 5	Product 5	15.99	150
324e6baa-75f5-4b34-a09d-f03de75433c5	Description of Product 3	Product 3	5.99	200
111723b3-2d0b-41b7-af20-e16f8ebf1079	Description of Product 4	Product 4	29.99	75

(5 rows)

3. Query to Retrieve All Suppliers and its output:

```
cqlsh:shopping_mall> SELECT * FROM suppliers;
```

supplier_id	address	email	name	phone
d5aa00d9-5642-4eb0-b6c9-59190a43f22b	789 Supplier St	supplier3@example.com	Supplier 3	789-012-3456
e4a6cd3f-09f4-4577-8822-668f61aa74af	012 Supplier St	supplier4@example.com	Supplier 4	012-345-6789
3a645acc-a98c-416f-a4af-e58c450150ea	456 Supplier St	supplier2@example.com	Supplier 2	456-789-0123
5b4beab8-712b-4276-a934-4f4d4c546255	123 Supplier St	supplier1@example.com	Supplier 1	123-456-7890
27c48f00-c5a4-49f1-882f-a66b0339791a	345 Supplier St	supplier5@example.com	Supplier 5	345-678-9012

(5 rows)